

Inflation for Classical Pair-Source Networks: Termination and Quantitative Obstructions

William Blair
willjblair1@gmail.com

September 2026

Abstract

The classical inflation hierarchy converges to the set of distributions compatible with a correlation scenario. We classify when convergence occurs at a finite order for pair-source scenarios with binary observations: a finite order characterizes compatibility if and only if every connected component of the source graph is a tree of diameter at most three. Order two then suffices. The classification holds for the Navascués–Wolfe hierarchy and its ancestral-independence and recursively expressible strengthenings; the latter two coincide on pair-source scenarios. On every other graph, each prescribed order accepts an incompatible, strictly positive rational law. For the triangle and square, the largest total-variation distance from compatibility among laws accepted at order t lies between c/t and $C/\sqrt{t}$. Corrected Fourier densities give the lower bound; a convex-order argument gives an upper bound that depends only on the number of observed outcomes. We also prove order-conversion bounds and certify finite-order separations between the hierarchies. Along a separate triangle family at distance d from compatibility, the first rejecting order is $\Theta(d^{-1/3})$. Over rational binary-triangle inputs of at most B bits, the worst required order is $2^{\Theta(B)}$. Exact rational certificates accompany the finite computations. The classification equivalence for binary observations and finite latent alphabets, together with several triangle results, is formalized in Lean.

Contents

1	Introduction	2
2	Pair-source scenarios and inflation	4
3	The termination classification	7
4	Quantitative convergence on the triangle and square	21
5	The defect family and rejecting-order growth	27
6	Rejecting orders from distance, input size and parity	35
7	Open problems	37
8	Verification and availability	38
A	Structural details of inflation	39
B	Order conversion and finite thresholds	41
C	Certificate formats and coverage	45
D	Supplementary inequalities and extraction results	47

1 Introduction

A correlation scenario has independent latent sources and observed variables that respond to the sources reaching them. Its compatible set need not be convex. Finite latent-cardinality bounds make this set semialgebraic, but no explicit characterization is known even for the classical triangle with binary observations [9, 20]. Inflation tests compatibility by copying each source t times and asking for a joint law on the copied observations with prescribed marginals [22]. Navascués and Wolfe proved asymptotic completeness of their hierarchy and asked whether, for a fixed scenario and fixed observed alphabets, some finite order already decides compatibility [15]. They identified the triangle as a possible counterexample.

We answer this question for pair-source scenarios, in which each source reaches exactly two observers. Such a scenario is a finite simple graph G : observers are vertices and sources are edges. We then study the order needed to reject incompatible laws near the compatible set.

1.1 Main results

Termination classification. For binary observations, a pair-source scenario has a finite characterizing order if and only if every connected component is a tree of diameter at most three, or *double-star*. In that case order two suffices. Otherwise, for every prescribed order there is an incompatible, strictly positive rational law accepted at that order (Theorem 3.2). This classification holds for the Navascués–Wolfe (NW), ancestral-independence (AI) and recursively expressible tests. The last two coincide on pair-source scenarios.

The proof reconstructs a model on each double-star. Every other connected graph contains an induced cycle or a five-vertex path; explicit witnesses on these subgraphs extend to the ambient graph. Figure 1 summarizes the classification.




Figure 1: The classification for binary pair-source scenarios. Vertices denote observers and edges independent pair sources. A graph terminates exactly when each component is a double-star. Here H is NW, AI or exp.

Quantitative convergence. Let $H_t^H(G)$ be the largest total-variation distance from compatibility among laws accepted at order t by hierarchy H . For the triangle and square we prove

$$\frac{c_G}{t} \leq H_t^H(G) \leq \frac{C_G}{\sqrt{t}},$$

with explicit positive constants, $c_{\square} = 1/47$ and $c_{\triangle} = 1/43$ (Corollary 4.9). The lower bounds use parity laws and corrected Fourier densities. The upper bound follows from a convex-order comparison and applies to any correlation scenario with K joint outcomes, with $C_G = \sqrt{K-1}/2$ (Theorem 4.8). On the square, increasing the NW order by a factor at most $3/2$ enforces all AI prescriptions (Theorem 4.10). Exact certificates bracket the parity thresholds of both tests through order ten and separate them at every even order in that range (Proposition 4.11).

Rejecting orders near the boundary. A second triangle construction places independent defects in a cube of copy indices. It gives passing witnesses for

$$Q(\varepsilon, r) = \varepsilon \delta_{000} + (1 - \varepsilon) \text{Bern}(r)^{\otimes 3}, \quad r \leq (1 - \varepsilon)^{t-1},$$

where $\text{Bern}(r)(1) = r$. For $P_{\varepsilon} = Q(\varepsilon, 1 - \frac{1}{2}\varepsilon^{2/3})$, a fan inequality gives a matching rejecting-order exponent:

$$t_{\min}^H(P_{\varepsilon}) = \Theta(\varepsilon^{-1/3}) = \Theta(d_{\text{TV}}(P_{\varepsilon}, \mathcal{C}_{\triangle})^{-1/3}).$$

We compute the leading distance constant and bound the rejecting-order constants (Proposition 5.13, Theorem 5.14). This exponent concerns this family alone; the worst-case bounds above cover all accepted laws. For rational binary-triangle inputs of at most B bits, the largest first rejecting order is $2^{\Theta(B)}$ (Proposition 6.2). A promised distance $\delta > 0$ gives the bound $\lfloor (K-1)/(4\delta^2) \rfloor + 1$ for any fixed correlation scenario (Corollary 6.1). On the square and the triangle, a law that violates one of two cubic parity certificates by v is rejected at order $\lfloor 3/v \rfloor + 1$ (Remark 6.3). The $1/n$ error estimate follows the approach of Navascués and Wolfe [15, Theorem 1]; the constant here does not depend on the number of sources.

1.2 Relation to prior work

Finite termination. Navascués and Wolfe prove order-two termination for stars and exhibit triangle-incompatible laws that pass order two [15, Section 4.1]. Our reconstruction extends their source-disjointness argument to double-stars, including their four-on-line example; the cycle and five-path witnesses complete the classification. Henson, Lal and Pusey [11] ask which causal structures are characterized by observed conditional independences. Their criterion and finite termination already differ on the four-on-line scenario, which fails the criterion but terminates at order two. The strict gap that Wolfe, Spekkens and Fritz [22] exhibit between expressible and AI prescriptions uses an observed node with observed children. Lemma 3.4 rules out that gap for the root-sink graphs considered here. The finite nonattainment example for polynomial optimization in [15] concerns a scenario in which every law is compatible; our obstruction concerns finite compatibility testing.

Inequalities and witnesses. The defect construction violates Finner’s inequality [6, 19]. The five-path obstruction uses the bilocal square-root inequality of Branciard, Rosset, Gisin and Pironio [4], with the endpoint observers serving as settings. Cycle incompatibility uses parity rigidity, developed for the triangle by Boreiri et al. [2, 3]; Renou and Beigi [17] prove token-counting rigidity on rings. We give the needed classical parity proofs for arbitrary latent spaces and stochastic responses. Our symmetric triangle parity family is among the families studied by da Silva, Pozas-Kerstjens and Parisio [5]. Inflation has also been used to certify nonlocality along continuous families [16]; our quantitative question is how the rejecting order grows as an incompatible law approaches the classical set.

The fan inequality follows the marginal-polytope method of [22, Section IV.D], also used for the wagon-wheel and web inequalities of Fraser and Wolfe [8]. We derive a formula indexed by the order and a second-moment bound. The witness densities and count-moment reductions use Fourier characters and Krawtchouk polynomials, standard tools in coding theory [14].

Convergence and effectivity. The convex-order estimate removes the source-count factor from the finite-sampling bound of [15]. Related convergence results for postselected inflation [10] and the combinatorial hierarchy [7] also depend on source parameters. The sampling comparison is related to Hoeffding’s convex-order inequality [12, Theorem 4]. The input-size upper bound combines our distance estimate with finite latent cardinality [20], real quantifier elimination [1] and a polynomial root bound. Quantum inflation and its convergent variants [21, 13] concern different compatible sets; their finite-termination questions lie outside the classification here.

Organization. Section 2 defines the scenarios and tests, with the triangle as an example. Section 3 proves the classification. Sections 4 and 5 give the convergence bounds and the defect-family analysis; Section 6 derives rejecting orders from distance, input size and parity certificates. Open problems and verification scope follow in Sections 7 and 8. The appendices contain the root-sink proof, order-conversion and threshold details, certificate formats and supplementary inequalities.

2 Pair-source scenarios and inflation

2.1 Scenarios

Definition 2.1 (Pair-source scenario). Let $G = (V, E)$ be a finite simple graph without isolated vertices. The associated scenario has one latent source S_e for each edge $e \in E$, taking values in an arbitrary measurable space, and one observed variable O_v with a finite alphabet $\mathcal{O}_v$ for each vertex $v \in V$. The sources are mutually independent, and

$$O_v \sim \kappa_v(\cdot \mid (S_e)_{e \ni v}),$$

conditionally independently across v given the sources. Write $\mathcal{C}_G$ for the set of joint laws of $(O_v)_{v \in V}$ obtained in this way. Unless stated otherwise, every $\mathcal{O}_v$ is $\{0, 1\}$; we then speak of binary observations.

Each source is read by exactly the two endpoints of its edge. We write C_m for the cycle on m vertices, P_n for the path on n vertices, and $\mathcal{C}_\Delta = \mathcal{C}_{C_3}$.

Sources shared by three or more observers, and latent variables with latent parents, are outside this definition; see Remark 3.9.

Assign each vertex v an independent random seed and append it to the source of an edge at v , which exists because no vertex is isolated; the other endpoint ignores the seed. The graph is unchanged and the seed makes the response at v deterministic, so stochastic and deterministic responses generate the same set $\mathcal{C}_G$ [22].

2.2 Inflation at order t

Definition 2.2 (Order- t inflation). Fix $t \geq 1$ and write $E(v) = \{e \in E : v \in e\}$ for the edges at v and $E(W) = \bigcup_{v \in W} E(v)$ for a vertex set W . The order- t inflation of G has independent copies

$S_e^1, \dots, S_e^t$ of each source, and for every vertex v and every index tuple $i \in [t]^{E(v)}$ one copied observation

$$O_v^i \sim \kappa_v(\cdot \mid (S_e^{i_e})_{e \in E(v)}),$$

all conditionally independent given the source copies. Writing r also for the constant tuple with value r , the *row* $r \in [t]$ is the family $(O_v^r)_{v \in V}$, and the t rows form the diagonal.

A vertex of degree d has t^d copies, so the inflation carries $\sum_v t^{\deg v}$ observed variables. The group $\prod_{e \in E} S_t$ acts on the copy labels: $\sigma = (\sigma_e)_{e \in E}$ sends the label i to $\sigma(i)$ with $\sigma(i)_e = \sigma_e(i_e)$, and hence sends O_v^i to $O_v^{\sigma(i)}$. Each σ maps rows to rows when all its components are equal, and otherwise mixes them.

Call a probability law on $\prod_{v \in V} (\mathcal{O}_v)^{[t]^{E(v)}}$ a *table* of order t . The three tests below ask for a table with prescribed marginals.

Definition 2.3 (Navascués–Wolfe feasible set). $P \in \mathcal{I}_t^{\text{NW}}(G)$ if there is a table Γ_t of order t such that

- (i) Γ_t is invariant under $\prod_{e \in E} S_t$; and
- (ii) its diagonal law is a tensor power,

$$\text{Law}_{\Gamma_t}((O_v^r)_{v \in V} : r = 1, \dots, t) = P^{\otimes t}.$$

Definition 2.4 (Injectable sets and the AI feasible set). A set S of copied observations is *injectable* if the map π that erases copy indices is injective on S and any two members of S whose vertices share an edge carry the same index on that edge [22, Definition 4]. Equivalently, S is injectable if and only if there are a vertex set $W \subseteq V$ and an assignment $\iota : E(W) \rightarrow [t]$ with

$$S = \{O_v^{\iota|_{E(v)}} : v \in W\},$$

in which case $\pi(S) = W$. Two sets of copied observations are *ancestrally independent* if the sets of source copies they read are disjoint. $P \in \mathcal{I}_t^{\text{AI}}(G)$ if $P \in \mathcal{I}_t^{\text{NW}}(G)$ with a witness Γ_t under which every injectable set S has marginal $P_{\pi(S)}$ and every union of pairwise ancestrally independent injectable sets has the product of the corresponding marginals.

For the equivalence in Definition 2.4, given S injectable, let $W = \pi(S)$ and let $\iota(e)$ be the common index carried on e by the members of S at the endpoints of e that lie in W , which is well defined by the agreement requirement; conversely every set of the displayed form is injectable. For the triangle this reproduces the description in Appendix A.3: the injectable sets are the subsets of the copied triangles.

Definition 2.5 (Recursively expressible feasible set). Expressible sets of copied observations are generated from injectable sets by the two rules of [22, Definition 7], read in the inflated causal graph of Definition 2.2: if X and Y are d -separated by Z and the laws on $X \cup Z$ and on $Y \cup Z$ are prescribed, then

$$p(x, y, z) = \frac{p(x, z) p(y, z)}{p(z)} \quad \text{where } p(z) > 0, \quad p(x, y, z) = 0 \quad \text{where } p(z) = 0,$$

is prescribed on $X \cup Y \cup Z$; and every marginal of a prescribed set is prescribed. $P \in \mathcal{I}_t^{\text{exp}}(G)$ if $P \in \mathcal{I}_t^{\text{AI}}(G)$ with a witness satisfying every prescription so generated.

Proposition 2.6 (Soundness). *For every pair-source scenario G and every $t \geq 1$,*

$$\mathcal{C}_G \subseteq \mathcal{I}_t^{\text{exp}}(G) \subseteq \mathcal{I}_t^{\text{AI}}(G) \subseteq \mathcal{I}_t^{\text{NW}}(G). \quad (1)$$

Proof. The two inclusions on the right hold because each test adds prescriptions to the previous one. For the first, realize P by a model, run the model on the inflated source family of Definition 2.2, and let Γ_t be the law of the copied observations. Permuting the copies of one source permutes independent identically distributed variables, so Γ_t is invariant. Distinct rows read disjoint source copies and each row has law P , so the diagonal law is $P^{\otimes t}$. An injectable set S reads, for each $v \in \pi(S)$, the copies indexed by a single global assignment ι , so the joint law of S is that of $(O_v)_{v \in \pi(S)}$ in the original model, namely $P_{\pi(S)}$. Ancestrally independent sets read disjoint source copies, hence are independent. The inflated model is a causal model for the inflated graph, so all conditional independences read off that graph hold, and the expressible prescriptions follow by induction on their generation. $\square$

For fixed P each prescription fixes a marginal of the table, so all three tests are linear feasibility problems. Restricting a table to fewer copy indices preserves the conditions, so the feasible sets decrease with t . The Navascués–Wolfe hierarchy is asymptotically complete, $\mathcal{I}_t^{\text{NW}}(G) \downarrow \mathcal{C}_G$ [15]. Consequently, for an incompatible law and $H \in \{\text{NW}, \text{AI}, \text{exp}\}$, the first rejecting order

$$t_{\min}^H(P) = \min\{t \geq 1 : P \notin \mathcal{I}_t^H(G)\}$$

is finite. We use total variation $d_{\text{TV}}(P, Q) = \frac{1}{2} \sum_x |P(x) - Q(x)|$ and the distances

$$d_{\text{TV}}(P, \mathcal{C}) = \inf_{Q \in \mathcal{C}} d_{\text{TV}}(P, Q), \quad d_2(P, \mathcal{C}) = \inf_{Q \in \mathcal{C}} \|P - Q\|_2.$$

The finite latent-cardinality theorem of Rosset, Gisin and Wolfe [20] makes $\mathcal{C}_G$ compact and semialgebraic for fixed finite observed alphabets.

2.3 The triangle as a worked example

The triangle has independent sources X, Y, Z and observations

$$A \sim \kappa_A(\cdot \mid X, Z), \quad B \sim \kappa_B(\cdot \mid X, Y), \quad C \sim \kappa_C(\cdot \mid Z, Y).$$

Its order- t inflation carries $3t^2$ observations, indexed as

$$A^{ij} = A(X_i, Z_j), \quad B^{ik} = B(X_i, Y_k), \quad C^{jk} = C(Z_j, Y_k), \quad i, j, k \in [t].$$

This notation allows stochastic responses as in Definition 2.2. A copied triangle is $\Delta_{ijk} = \{A^{ij}, B^{ik}, C^{jk}\}$, with source copies X_i, Z_j, Y_k ; the diagonal triangles are $\Delta_{\ell\ell\ell}$ for $\ell \in [t]$.

We omit the graph argument from $\mathcal{I}_t^{\text{NW}}$, $\mathcal{I}_t^{\text{AI}}$ and $\mathcal{I}_t^{\text{exp}}$ when discussing the triangle. Definition 2.3 is the full-table symmetric formulation of [15, equations (16)–(19)]; its lower-degree diagonal conditions follow by marginalization. Their cyclic indexing and ours describe the same copied variables. The injectable sets are exactly the subsets of copied triangles (Appendix A.3). The AI conditions also prescribe products for sets with disjoint copied ancestors, for example

$$\text{Law}(A^{12}, B^{21}, C^{12}) = P_A \otimes P_B \otimes P_C.$$

The NW test does not impose this relation at order two.

Every compatible triangle law satisfies the event form of Finner’s inequality,

$$P(000)^2 \leq P_A(0)P_B(0)P_C(0). \quad (2)$$

Lemma 5.7 proves it for arbitrary latent spaces and stochastic responses [6, 19].




Figure 2: Sources (squares) and observations (circles) in the triangle. Each copied observation carries the indices of its two source parents. Diagonal triangles $\Delta_{\ell\ell\ell}$ use disjoint source copies; the NW test requires their joint law to be $P^{\otimes t}$. The full order- t table includes all $3t^2$ copied observations.

3 The termination classification

Definition 3.1 (Double-star). A *double-star* is a graph with a distinguished edge $\{b, c\}$, the *central edge*, such that every other vertex is adjacent to exactly one of b and c and to nothing else. Equivalently, a double-star is a tree of diameter at most three. Single edges and stars are double-stars.

Theorem 3.2 (Termination classification). *Let G be a finite simple graph without isolated vertices, with binary observations. The following are equivalent.*

- (a) $\mathcal{I}_t^{\text{NW}}(G) = \mathcal{C}_G$ for some finite t .
- (b) $\mathcal{I}_t^{\text{AI}}(G) = \mathcal{C}_G$ for some finite t .
- (c) $\mathcal{I}_t^{\text{exp}}(G) = \mathcal{C}_G$ for some finite t .
- (d) Every connected component of G is a double-star.

When they hold,

$$\mathcal{I}_2^{\text{NW}}(G) = \mathcal{I}_2^{\text{AI}}(G) = \mathcal{I}_2^{\text{exp}}(G) = \mathcal{C}_G.$$

When they fail, for every $t \geq 1$ there is a strictly positive rational law

$$P \in \mathcal{I}_t^{\text{exp}}(G) \setminus \mathcal{C}_G \subseteq \mathcal{I}_t^{\text{AI}}(G) \setminus \mathcal{C}_G \subseteq \mathcal{I}_t^{\text{NW}}(G) \setminus \mathcal{C}_G.$$

The graph and the alphabets stay fixed as t grows; only the law P varies. The proof combines Theorem 3.10 with Theorems 3.21 and 3.26 through the exhaustion lemma below.

3.1 Structural lemmas

In the inflated graph every source copy is a root and every copied observation is a sink, and as a result the AI and expressible prescriptions coincide; Appendix A.1 proves the equivalence.

Definition 3.3 (Shared-parent graph and AI sets). For a set W of copied observations, the *shared-parent graph* on W joins two members when they read a common source copy. Call W an *AI set* if it is a union of pairwise ancestrally independent injectable sets.

Lemma 3.4 (Root-sink lemma). *Let G be a pair-source scenario and $t \geq 1$.*

- (i) *A set W of copied observations is an AI set if and only if every connected component of its shared-parent graph is injectable.*
- (ii) *A table satisfies every expressible prescription of Definition 2.5 if and only if it satisfies the injectable and ancestral-independence prescriptions of Definition 2.4.*

In particular $\mathcal{I}_t^{\text{exp}}(G) = \mathcal{I}_t^{\text{AI}}(G)$ for every t .

Lemma 3.4 identifies the two stronger tests. It says nothing about the Navascués–Wolfe test and imposes no causal model on the table as a whole.

Source-disjoint blocks.

Lemma 3.5 (Source-disjoint independence). *Let $t \geq 2$ and $P \in \mathcal{I}_t^{\text{NW}}(G)$. If $U, W \subseteq V$ satisfy $E(U) \cap E(W) = \emptyset$, then*

$$P_{U \cup W} = P_U \otimes P_W.$$

Proof. Let Γ be a witness. Let σ be the element of $\prod_e S_t$ that transposes the copy indices 1 and 2 on every edge of $E(W)$ and fixes all other indices. For $u \in U$ every edge at u lies outside $E(W)$, so σ fixes O_u^1 ; for $w \in W$ every edge at w lies in $E(W)$, so σ sends O_w^1 to O_w^2 . Invariance of Γ gives

$$\text{Law}((O_u^1)_{u \in U}, (O_w^1)_{w \in W}) = \text{Law}((O_u^1)_{u \in U}, (O_w^2)_{w \in W}).$$

The left side is the marginal of row 1, namely $P_{U \cup W}$. The right side reads U -coordinates from row 1 and W -coordinates from row 2, and the two rows are independent with law P each, so it is $P_U \otimes P_W$. $\square$

Two vertices joined by an edge have that edge in both incidence sets, so the hypothesis says exactly that U and W are disjoint and no edge joins them. Iterating the lemma with U a single block and W the union of the others gives mutual independence of any family of pairwise source-disjoint blocks. In particular the observed law of a disconnected graph factors over its components from order two on.

Induced subgraphs.

Lemma 3.6 (Induced-subgraph transport). *Let H be an induced subgraph of G in which no vertex is isolated, and let $\mathcal{J} \in \{\mathcal{I}^{\text{NW}}, \mathcal{I}^{\text{AI}}, \mathcal{I}^{\text{exp}}\}$.*

- (i) *If $Q \in \mathcal{C}_G$ then $Q_{V(H)} \in \mathcal{C}_H$.*
- (ii) *If $P_H \in \mathcal{J}_t(H)$ then $P := P_H \otimes \text{Bern}(1/2)^{\otimes (V \setminus V(H))} \in \mathcal{J}_t(G)$.*
- (iii) *$d_{\text{TV}}(P, \mathcal{C}_G) \geq d_{\text{TV}}(P_H, \mathcal{C}_H)$. In particular, if $P_H \notin \mathcal{C}_H$ then $P \notin \mathcal{C}_G$.*

The proof is in Appendix A.2.

Part (ii) preserves strict positivity and rationality of the target, since the appended factor is the fair-bit law.

3.2 Reduction to cycles and paths

Lemma 3.7 (Exhaustion). *Let K be a finite connected graph that has at least one edge and is not a double-star. Then K has an induced subgraph isomorphic to a cycle C_m with $m \geq 3$ or to the path P_5 .*

Proof. Suppose first that K contains a cycle, and let C be one of minimum length. A chord of C would split it into two shorter cycles, so C has no chords and is induced. Its length is at least 3.

Suppose next that K is a tree. We claim that a tree of diameter at most three is a double-star. If the diameter is 1 then K is a single edge. If it is 2 then K is a star. Let the diameter be 3 and fix a geodesic $u_0 u_1 u_2 u_3$. Let w be a vertex other than u_1 and u_2 , and consider the unique path from w to u_1 . If it avoids u_2 , then it also avoids u_3 , since the path from u_3 to u_1 passes through u_2 ; so concatenating with $u_1 u_2 u_3$ gives a path and $d(w, u_3) = d(w, u_1) + 2 \leq 3$, whence $d(w, u_1) = 1$. If it contains u_2 , then the path from w to u_2 avoids u_1 and hence u_0 , and the same argument with u_0 gives $d(w, u_2) = 1$. Every vertex is therefore adjacent to u_1 or u_2 , and K is a double-star with central edge $u_1 u_2$.

A tree that is not a double-star therefore has diameter at least 4. Take a geodesic of length at least 4 and let $v_0, \dots, v_4$ be five consecutive vertices on it. Any edge of K among them other than the four geodesic edges would create a cycle in a tree, so the induced subgraph on $\{v_0, \dots, v_4\}$ is P_5 . $\square$

Fritz classified the graph scenarios in which every correlation is classical as the disjoint unions of stars [9, Theorem 3.8]: every other connected graph has an induced triangle, square or four-vertex path [9, Lemma 3.9], and a correlation on that induced subgraph extends to the whole graph by deterministic outcomes at the remaining vertices. Lemma 3.7 and the proof below follow that argument, with double-stars in place of stars, cycles and the five-vertex path as the forbidden induced subgraphs, and inflation witnesses in place of correlations. Fritz also framed scenarios with arbitrary sources as hypergraphs [9, Section 3].

Proof of Theorem 3.2. Assume (d). Each component is a double-star with binary observations, so Theorem 3.10 and Corollary 3.13 give $\mathcal{I}_2^{\text{NW}}(G) = \mathcal{C}_G$. With Proposition 2.6,

$$\mathcal{C}_G \subseteq \mathcal{I}_2^{\text{exp}}(G) \subseteq \mathcal{I}_2^{\text{AI}}(G) \subseteq \mathcal{I}_2^{\text{NW}}(G) = \mathcal{C}_G,$$

so all four sets coincide and (a), (b), (c) hold at $t = 2$.

Assume (d) fails, and fix $t \geq 1$. Some component K of G is not a double-star, and K has an edge because G has no isolated vertices. By Lemma 3.7, K has an induced subgraph H isomorphic to C_m for some $m \geq 3$ or to P_5 ; since K is induced in G , so is H , and H has no isolated vertices. Theorem 3.21 in the first case and Theorem 3.26 in the second supply a strictly positive rational law

$$P_H \in \mathcal{I}_t^{\text{exp}}(H) \setminus \mathcal{C}_H.$$

Let $P = P_H \otimes \text{Bern}(1/2)^{\otimes(V \setminus V(H))}$. By Lemma 3.6(ii), $P \in \mathcal{I}_t^{\text{exp}}(G)$, and by Lemma 3.6(iii), $P \notin \mathcal{C}_G$. It is strictly positive and rational because P_H is. By Proposition 2.6 the same law lies in $\mathcal{I}_t^{\text{AI}}(G)$ and in $\mathcal{I}_t^{\text{NW}}(G)$, so none of the three tests equals $\mathcal{C}_G$ at order t . As t was arbitrary, (a), (b) and (c) all fail.

This proves the equivalence, the order-two bound and the witness assertion. $\square$

For $m = 3$ the witness of Theorem 3.21 is a second nonterminating family for the triangle, alongside the defect construction of Theorem 5.2.

Remark 3.8 (Larger observed alphabets). Fix finite alphabets $\mathcal{O}_v$ with $|\mathcal{O}_v| \geq 2$. The classification is unchanged, with a different order in the terminating case. If some component is not a double-star, regard the binary law P of Theorem 3.2 as a law on $\prod_v \mathcal{O}_v$ supported on $\{0, 1\}^V$. A compatible realization of such a law has responses that leave $\{0, 1\}$ with probability zero, hence is a binary realization, so the law remains incompatible; and its witness, read in the larger alphabets, satisfies the same prescriptions. The transferred law is supported on the binary sub-alphabet, so the strict positivity in Theorem 3.2 is a statement about binary observations. If every component is a double-star, Theorem 3.10 gives termination at

$$T = \max(\{2\} \cup \{|\mathcal{O}_v| : v \text{ a leaf}\}),$$

the leaves being taken with respect to a choice of central edge in each component; the bound does not involve the alphabets at the central vertices.

Remark 3.9 (Scenarios that are not pair-source). Let three observers read one common source. That source can carry a sample of the whole observed triple, so every law on the three outcomes is compatible [9, Proposition 3.7] and the order-one test characterizes the compatible set. Its observer adjacency graph is the same as that of the pair-source triangle, whose hierarchy does not terminate. Thus the adjacency graph alone does not decide termination when a source may be read by three or more observers.

3.3 Finite reconstruction on double-stars

Fix a double-star G in the sense of Definition 3.1: two central vertices b, c joined by the central edge $y = \{b, c\}$, left leaves $\ell_1, \dots, \ell_p$ joined to b by edges $e_1, \dots, e_p$, and right leaves $r_1, \dots, r_q$ joined to c by edges $f_1, \dots, f_q$. Observed alphabets $\mathcal{O}_v$ are finite and arbitrary. Put

$$T = \max(\{2\} \cup \{|\mathcal{O}_{\ell_s}|\}_{s \leq p} \cup \{|\mathcal{O}_{r_j}|\}_{j \leq q}).$$

In the order- T inflation the leaves have one index each, and the centres have one index per incident edge. Write

$$L_s^m = O_{\ell_s}^m, \quad R_j^m = O_{r_j}^m, \quad B^{\mathbf{i}, n} = O_b^{(\mathbf{i}, n)}, \quad C^{n, \mathbf{k}} = O_c^{(n, \mathbf{k})},$$

where $m, n \in [T]$, $\mathbf{i} \in [T]^p$ records the index on each e_s , and $\mathbf{k} \in [T]^q$ the index on each f_j . Row m consists of L_s^m , R_j^m , $B^{m, m}$ and $C^{m, m}$, with $\mathbf{i}$ and $\mathbf{k}$ constant equal to m .

For a law P on $\prod_v \mathcal{O}_v$ write $P(\beta, \gamma, \xi, \zeta)$ for the probability that $O_b = \beta$, $O_c = \gamma$, $O_{\ell_s} = \xi_s$ for all s and $O_{r_j} = \zeta_j$ for all j .

Theorem 3.10 (Double-star reconstruction). *For every double-star G with finite observed alphabets,*

$$\mathcal{I}_T^{\text{NW}}(G) = \mathcal{C}_G.$$

Proof. The inclusion $\mathcal{C}_G \subseteq \mathcal{I}_T^{\text{NW}}(G)$ is Proposition 2.6. Let $P \in \mathcal{I}_T^{\text{NW}}(G)$ and let Γ be a witness at order T . We recover a model for P in four steps.

(i) *The leaves and all their copies are independent.* Distinct leaves have disjoint incidence sets, each consisting of a single edge, so Lemma 3.5, available because $T \geq 2$, applies to one leaf against the union of the others and, iterated, gives

$$P_{\text{leaf}}(\xi, \zeta) := P(O_{\ell_s} = \xi_s \ \forall s, \ O_{r_j} = \zeta_j \ \forall j) = \prod_s P_{\ell_s}(\xi_s) \prod_j P_{r_j}(\zeta_j). \quad (3)$$

Under Γ the T rows are independent with law P each, and the leaf entries of a single row are independent by (3). Every copied leaf observation occurs in exactly one row, so the whole family $\{L_s^m, R_j^m\}$ is independent under Γ , with $L_s^m \sim P_{\ell_s}$ and $R_j^m \sim P_{r_j}$.

(ii) *A positive conditioning event.* For each left leaf choose $x_s^1, \dots, x_s^T \in \mathcal{O}_{\ell_s}$ so that every symbol of $\text{supp } P_{\ell_s}$ occurs in the list, which is possible because $T \geq |\mathcal{O}_{\ell_s}|$; fill the remaining positions with repetitions of a supported symbol. Choose $z_j^1, \dots, z_j^T$ for each right leaf in the same way. Let

$$E = \{L_s^m = x_s^m \ \forall s, m\} \cap \{R_j^m = z_j^m \ \forall j, m\}.$$

By (i),

$$\Gamma(E) = \prod_{s, m} P_{\ell_s}(x_s^m) \prod_{j, m} P_{r_j}(z_j^m) > 0. \quad (4)$$

(iii) *The conditional marginal of the two centres.* Fix $\mathbf{i} \in [T]^p$ and $\mathbf{k} \in [T]^q$. Choose permutations $\sigma_s \in S_T$ of the copy indices of e_s with $\sigma_s(1) = i_s$, permutations $\tau_j \in S_T$ of those of f_j with $\tau_j(1) = k_j$, and the identity on the central edge y . The resulting element of the symmetry group carries row m to

$$(B^{(\sigma_s(m))_s, m}, C^{m, (\tau_j(m))_j}, (L_s^{\sigma_s(m)})_s, (R_j^{\tau_j(m)})_j),$$

which for $m = 1$ is $(B^{\mathbf{i}, 1}, C^{1, \mathbf{k}}, (L_s^{i_s})_s, (R_j^{k_j})_j)$. By invariance the joint law of these T images equals the joint law of the T rows, namely $P^{\otimes T}$. As m runs over $[T]$ the index $\sigma_s(m)$ runs over $[T]$, so every copied leaf observation occurs exactly once among the images and the event E is determined by them. Marginalizing the images of rows $2, \dots, T$ over their two centre entries and using (3),

$$\Gamma(B^{\mathbf{i}, 1} = \beta, C^{1, \mathbf{k}} = \gamma, E) = P(\beta, \gamma, x^{\mathbf{i}}, z^{\mathbf{k}}) \prod_s \prod_{m \neq i_s} P_{\ell_s}(x_s^m) \prod_j \prod_{m \neq k_j} P_{r_j}(z_j^m),$$

where $x^{\mathbf{i}} = (x_s^{i_s})_s$ and $z^{\mathbf{k}} = (z_j^{k_j})_j$. Dividing by (4) and using (3) once more,

$$\Gamma(B^{\mathbf{i}, 1} = \beta, C^{1, \mathbf{k}} = \gamma \mid E) = P(O_b = \beta, O_c = \gamma \mid O_{\ell_s} = x_s^{i_s} \forall s, O_{r_j} = z_j^{k_j} \forall j), \quad (5)$$

the conditioning event on the right having positive probability whenever every $x_s^{i_s}$ and $z_j^{k_j}$ is supported. This identity uses only the symmetry and the diagonal law of the witness.

(iv) *Two response tables and fresh sources.* For $\xi \in \prod_s \text{supp } P_{\ell_s}$ pick positions $\mathbf{i}(\xi)$ with $x_s^{i_s(\xi)} = \xi_s$, and for $\zeta \in \prod_j \text{supp } P_{r_j}$ pick $\mathbf{k}(\zeta)$ with $z_j^{k_j(\zeta)} = \zeta_j$. Under the conditional law $\Gamma(\cdot \mid E)$ define the random maps

$$F(\xi) = B^{\mathbf{i}(\xi), 1}, \quad G(\zeta) = C^{1, \mathbf{k}(\zeta)},$$

with arbitrary fixed values on unsupported arguments. Both are functions of the same table, so the pair (F, G) has a joint law μ , and (5) reads

$$\mu(F(\xi) = \beta, G(\zeta) = \gamma) = P(O_b = \beta, O_c = \gamma \mid \xi, \zeta) \quad (6)$$

for all supported ξ, ζ .

Now take fresh independent sources: let the central source be $Y^* = (F, G) \sim \mu$, let the source of e_s be $X_s^* \sim P_{\ell_s}$ and the source of f_j be $Z_j^* \sim P_{r_j}$, all mutually independent. Respond by

$$O_{\ell_s} = X_s^*, \quad O_{r_j} = Z_j^*, \quad O_b = F(X_1^*, \dots, X_p^*), \quad O_c = G(Z_1^*, \dots, Z_q^*).$$

Each response reads only the sources of its incident edges: b reads $X_1^*, \dots, X_p^*$ and Y^* , and c reads Y^* and $Z_1^*, \dots, Z_q^*$. This is a model for the scenario G , and by (6) and (3) its observed law is

$$\prod_s P_{\ell_s}(\xi_s) \prod_j P_{r_j}(\zeta_j) \mu(F(\xi) = \beta, G(\zeta) = \gamma) = P_{\text{leaf}}(\xi, \zeta) P(\beta, \gamma \mid \xi, \zeta) = P(\beta, \gamma, \xi, \zeta).$$

The model never samples unsupported leaf values, so the arbitrary choices there do not matter. Hence $P \in \mathcal{C}_G$. $\square$

A central source with $|\mathcal{O}_b| \prod_s |\text{supp } P_{\ell_s}| |\mathcal{O}_c| \prod_j |\text{supp } P_{r_j}|$ values suffices: a left response table is a map on the product of the left leaf supports, and likewise on the right. Empty products equal one.

Corollary 3.11 (Binary observations). *For a double-star with binary observations, $\mathcal{I}_2^{\text{NW}}(G) = \mathcal{C}_G$.*

Proof. All leaf alphabets have two elements, so $T = 2$. $\square$

Corollary 3.12 (Four-observer path). *Consider $A(X)$, $B(X, Y)$, $C(Y, Z)$, $D(Z)$, the path P_4 with A and D at the ends. Then $\mathcal{I}_T^{\text{NW}}(P_4) = \mathcal{C}_{P_4}$ with $T = \max\{2, |\mathcal{O}_A|, |\mathcal{O}_D|\}$. For binary endpoints $T = 2$, whatever the finite alphabets of B and C .*

Proof. P_4 is the double-star with central edge $\{B, C\}$, one left leaf A and one right leaf D ; apply Theorem 3.10, whose bound involves only the leaf alphabets. $\square$

Fritz [9, Theorem 2.4] identifies classical correlations on P_4 with Bell-local conditional distributions, with the endpoint observations serving as settings. The four-observer path is the four-on-line scenario of [15, Figure 8]. Navascués and Wolfe list it with the triangle among the scenarios whose compatible set is cut out by inequality constraints, show that order two does not suffice for the triangle, and leave the four-on-line case open. The corollary settles it at order two for binary endpoints.

Corollary 3.13 (Disjoint unions). *Let G be a disjoint union of double-stars $G_1, \dots, G_n$ with orders $T_1, \dots, T_n$ as above, and let $T = \max\{2, T_1, \dots, T_n\}$. Then $\mathcal{I}_T^{\text{NW}}(G) = \mathcal{C}_G$.*

Proof. Let $P \in \mathcal{I}_T^{\text{NW}}(G)$ with witness Γ . Distinct components are source-disjoint, so Lemma 3.5, iterated, gives $P = \bigotimes_a P_{V(G_a)}$. The restriction of Γ to the copied observations at the vertices of G_a is a table for the order- T inflation of G_a : it is invariant under permutations of the copy indices of the edges of G_a , and its diagonal law is the corresponding marginal of $P^{\otimes T}$, namely $(P_{V(G_a)})^{\otimes T}$. Hence $P_{V(G_a)} \in \mathcal{I}_T^{\text{NW}}(G_a) \subseteq \mathcal{I}_{T_a}^{\text{NW}}(G_a)$, the inclusion because discarding the copy indices outside $[T_a]$ leaves a witness at order T_a . Theorem 3.10 now gives $P_{V(G_a)} \in \mathcal{C}_{G_a}$. Realize each component by its own model on its own independent source family; the joint observed law is the product, which is P . $\square$

A star is the case $q = 0$, with c one of its leaves. There the bound of Theorem 3.10 is not sharp: order two suffices with arbitrary finite alphabets [15, Section 4.1], by taking independent sources equal to the leaf values and letting the centre respond with its conditional law given them. Theorem 3.10 adds the central edge, across which the two centres share a source, so the leaf values on the two sides have to be assembled into one pair of response tables.

3.4 Cycles

For every cycle and every order, we construct a strictly positive law that passes all three tests and has positive distance from compatibility.

Throughout, the m -cycle C_m , $m \geq 3$, has observed variables $O_0, \dots, O_{m-1}$ and sources $S_0, \dots, S_{m-1}$, indices read modulo m , with O_v a child of S_{v-1} and S_v ; the copied observations of Definition 2.2 are written O_v^{ij} , children of $S_{v-1,i}$ and $S_{v,j}$. The triangle is C_3 , with $\mathcal{C}_\Delta = \mathcal{C}_{C_3}$, read as $A = O_0$, $B = O_1$, $C = O_2$ and $X = S_0$, $Y = S_1$, $Z = S_2$; the square is C_4 , written $\square$, read as $A(X, W)$, $B(X, Y)$, $C(Y, Z)$, $D(Z, W)$ with $X = S_0$, $Y = S_1$, $Z = S_2$, $W = S_3$.

3.4.1 Parity witnesses on every cycle

Lemma 3.14 (Auxiliary parity density). *For $N \geq 1$ and $q \geq 0$ let*

$$H_{N,q}(s) = 2^{-N} \sum_{\substack{S \subseteq [N] \\ |S| \text{ even}}} (-q)^{|S|/2} \prod_{w \in S} s_w, \quad s \in \{-1, 1\}^N. \quad (7)$$

Then $H_{N,q}$ is normalized, its characters are

$$\mathbb{E}_{H_{N,q}} \left[\prod_{w \in S} s_w \right] = \begin{cases} 0, & |S| \text{ odd}, \\ (-q)^{|S|/2}, & |S| \text{ even}, \end{cases} \quad (8)$$

and if $N^2 q \leq 1/4$ then $H_{N,q}(s) \geq \frac{2}{3} 2^{-N} > 0$ for every s .

Proof. Normalization and (8) are orthogonality of the characters of $\{-1, 1\}^N$ under the uniform measure. For positivity, $\binom{N}{2j} \leq N^{2j}$, so

$$|2^N H_{N,q}(s) - 1| \leq \sum_{j \geq 1} \binom{N}{2j} q^j \leq \sum_{j \geq 1} (N^2 q)^j \leq \sum_{j \geq 1} 4^{-j} = \frac{1}{3}. \quad \square$$

For $F \subseteq \{0, \dots, m-1\}$ let $\partial F = \{w : |F \cap \{w, w+1\}| = 1\}$ be the set of edges of C_m incident to an odd number of vertices of F . Every vertex meets two edges, so $|\partial F|$ is even.

Definition 3.15 (Cycle parity family). For $m \geq 3$ and $q \geq 0$ let $P_{m,q}$ be the signed measure on $\{-1, 1\}^m$ with density $2^{-m} \sum_F (-q)^{|\partial F|/2} \prod_{v \in F} O_v$ against the counting measure, so that it has total mass 1 and

$$\mathbb{E}_{P_{m,q}} \left[\prod_{v \in F} O_v \right] = (-q)^{|\partial F|/2} \quad \text{for every } F. \quad (9)$$

If $m^2 q \leq 1/4$ then $P_{m,q}$ is a probability law, namely the law of $(\sigma_{v-1} \sigma_v)_{v=0}^{m-1}$ for $\sigma \sim H_{m,q}$. In the two cases we use, the exact ranges are wider.

Taking F the whole vertex set gives $\mathbb{E}[\prod_v O_v] = 1$, so wherever $P_{m,q}$ is a probability law the product of all m signs is 1 almost surely; taking F a contiguous arc gives mean $-q$. For $m = 3$ the family is the parity-even law with observed means x, y, z ,

$$\Pi(x, y, z)(\alpha, \beta, \gamma) = \frac{1}{4}(1 + x\alpha + y\beta + z\gamma) \quad \text{on } \alpha\beta\gamma = 1, \quad P_{3,q} = \Pi(-q, -q, -q), \quad (10)$$

which is a probability law exactly for $0 \leq q \leq 1/3$; and for $m = 4$ it is the square parity family $P_q := P_{4,q}$ with atoms

$$\frac{1}{8}(1-6q+q^2) \text{ at } 0000, \quad \frac{1}{8}(1-q^2) \text{ at } 0011, 0110, 1001, 1100, \quad \frac{1}{8}(1+q)^2 \text{ at } 0101, 1010, 1111, \quad (11)$$

zero on odd parity, in the convention $0 \leftrightarrow +1$. It is a law exactly for $0 \leq q \leq 3 - 2\sqrt{2}$.

Fix $t \geq 1$, put $N = mt$ and draw auxiliary signs $s = (s_{v,i})_{v \in \{0, \dots, m-1\}, i \in [t]}$ from $H_{N,q}$. Define every copied observation by

$$O_v^{ij} = s_{v-1,i} s_{v,j}. \quad (12)$$

For a set F of copied observations let ∂F denote the set of copied sources occurring an odd number of times in $\prod_{O \in F} O$; each copied observation contributes two incidences, so $|\partial F|$ is even, and (8) gives

$$\mathbb{E} \left[\prod_{O \in F} O \right] = (-q)^{|\partial F|/2}. \quad (13)$$

Lemma 3.16 (Boundaries of prescribed characters). *Let $F_1, \dots, F_r$ be pairwise ancestrally independent injectable sets of copied observations of C_m and let $U_i \subseteq F_i$. Then each ∂U_i has even cardinality and contains at most one copy of each source, so it is empty or meets at least two source families; the sets $\partial U_1, \dots, \partial U_r$ are pairwise disjoint; and $\partial(U_1 \cup \dots \cup U_r) = \bigsqcup_i \partial U_i$, which again is empty or meets at least two source families. Moreover, erasing copy indices maps ∂U_i bijectively onto the boundary of the corresponding original observed set.*

Proof. An injectable set contains at most one copied observation over each original vertex and its ancestry contains at most one copy of each original source, so ∂U_i contains at most one coordinate from each family; being of even cardinality it is empty or has at least two elements in distinct families. Erasing copy indices carries the incidence count of U_i to that of its image, which gives the last claim. Ancestrally independent sets have disjoint ancestries, hence disjoint coordinate sets and disjoint boundaries, and the boundary of a union of sets with disjoint ancestries is the disjoint union of the boundaries. If that union is nonempty then some ∂U_i is nonempty and already meets two families. $\square$

Lemma 3.17 (The parity witness passes every prescription). *Let $m \geq 3$, $t \geq 1$ and $0 < q \leq 1/(4m^2t^2)$. The auxiliary density is strictly positive, and the law of (12) is a witness placing $P_{m,q}$ in $\mathcal{I}_t^{\text{AI}}(C_m) = \mathcal{I}_t^{\text{exp}}(C_m)$.*

Proof. Positivity of $H_{N,q}$ holds since $N^2q = m^2t^2q \leq 1/4$. The density (7) is symmetric in the N coordinates, so the witness is invariant under independent permutations of the copy indices of each source. By Lemma 3.16 an injectable set F has ∂F equal to the boundary of its image, so (13) and (9) give it the prescribed marginal of $P_{m,q}$. For pairwise ancestrally independent injectable $F_1, \dots, F_r$ and any $U_i \subseteq F_i$, Lemma 3.16 and (13) give

$$\mathbb{E}\left[\prod_i \prod_{O \in U_i} O\right] = (-q)^{\frac{1}{2} \sum_i |\partial U_i|} = \prod_i \mathbb{E}\left[\prod_{O \in U_i} O\right],$$

and characters determine a law on a finite sign cube, so the blocks are independent with the prescribed marginals. The diagonal rows are pairwise ancestrally independent injectable copied cycles, so the diagonal law is $P_{m,q}^{\otimes t}$. Lemma 3.4 supplies the recursive prescriptions. $\square$

The witness is a law on the copied observations. It does not come from an inflated model, since under $H_{N,q}$ the auxiliary signs are correlated across families.

Lemma 3.18 (Exact parity rigidity). *Let $R \in \mathcal{C}_\Delta$ with $ABC = 1$ almost surely. Then there are $\sigma, \tau, v \in [-1, 1]$ with $(\mathbb{E}A, \mathbb{E}B, \mathbb{E}C) = (\sigma\tau, \sigma v, \tau v)$. In particular $\mathbb{E}A \cdot \mathbb{E}B \cdot \mathbb{E}C \geq 0$.*

Proof. Absorb the response seeds into incident sources, so that $A = A(X, Z)$, $B = B(X, Y)$, $C = C(Z, Y)$ are deterministic signs. For μ_Y -almost every y the identity $A(X, Z)B(X, y)C(Z, y) = 1$ holds for $\mu_X \otimes \mu_Z$ -almost every (X, Z) ; fix such a y_0 and put $S(X) = B(X, y_0)$, $T(Z) = C(Z, y_0)$, so that $A = ST$ almost surely. Substituting in $ABC = 1$ gives $(S(X)B(X, Y))(T(Z)C(Z, Y)) = 1$ almost surely. Write $B' = SB$ and $C' = TC$. Conditionally on Y these are independent signs, and they agree almost surely, so each equals a function $U(Y)$ almost surely. Hence $A = ST$, $B = SU$, $C = TU$ almost surely, and independence of the three sources gives the means with $\sigma = \mathbb{E}S$, $\tau = \mathbb{E}T$, $v = \mathbb{E}U$. Their product is $(\sigma\tau v)^2$. $\square$

Exact rigidity of the parity-perfect (parity token-counting) triangle correlations is due to Boreiri, Girardin, Ulu, Lipka-Bartosik, Brunner and Sekatski [2], and Boreiri, Ulu, Brunner

and Sekatski [3] give a noise-robust form (a compatible law with parity error ϵ is within 3ϵ of a parity-perfect compatible law); Lemma 3.19 is a moment form of that robustness. Renou and Beigi [17, 18] prove rigidity of token-counting distributions on every ring scenario with bipartite sources; these are the prior rigidity results on cycles, but they concern integer token counts rather than parity and do not contain the lemma below. We include both proofs: we need the statements for arbitrary latent alphabets and stochastic responses, and the constant of Lemma 3.19 enters the distance bounds below.

Lemma 3.19 (Quantitative parity rigidity). *Let $R \in \mathcal{C}_\Delta$ and $W = \mathbb{E}_R[ABC]$. Then*

$$\max\{\mathbb{E}_R A, \mathbb{E}_R B, \mathbb{E}_R C\} \geq -4(1 - W).$$

Proof. Let $a(x, z)$, $b(x, y)$, $c(z, y)$ be the conditional means of A, B, C given the sources; they take values in $[-1, 1]$, and since the three responses use independent private randomness, $W = \mathbb{E}[abc]$ and $\mathbb{E}_R A = \mathbb{E}[a]$, with the expectations over the independent sources X, Y, Z . We use two inequalities: for $p, r \in [-1, 1]$,

$$|p - r| \leq 1 - pr, \quad (14)$$

since for $p \geq r$ the difference $(1 - pr) - (p - r) = (1 - p)(1 + r)$ is nonnegative, and for $u, w \in [-1, 1]$,

$$|u + w| \geq 2uw, \quad (15)$$

since for $uw \leq 0$ the right side is nonpositive and for u, w of the same sign $|u| + |w| - 2|u||w| = |u|(1 - |w|) + |w|(1 - |u|) \geq 0$.

Choose y_0 with $\mathbb{E}_{X,Z}[1 - a(X, Z)b(X, y_0)c(Z, y_0)] \leq 1 - W$; it exists because the average over Y of the left side is $1 - W$. Put $S(X) = b(X, y_0)$ and $T(Z) = c(Z, y_0)$. By (14) with $p = a$, $r = ST$,

$$\mathbb{E}[a - ST] \leq \mathbb{E}[1 - aST] \leq 1 - W, \quad \text{so} \quad |\mathbb{E}[a] - \mathbb{E}[S] \mathbb{E}[T]| \leq 1 - W,$$

using the independence of X and Z . Next let $u(y) = \mathbb{E}_X[S(X)b(X, y)]$ and $w(y) = \mathbb{E}_Z[T(Z)c(Z, y)]$, both in $[-1, 1]$, and $v(y) = \text{sgn}(u(y) + w(y))$. By (15),

$$\begin{aligned} \mathbb{E}_Y[1 - vu] + \mathbb{E}_Y[1 - vw] &= \mathbb{E}_Y[2 - |u + w|] \leq 2 - 2\mathbb{E}_Y[uw] \\ &= 2 - 2\mathbb{E}[STbc] \leq 2 - 2(W - \mathbb{E}[a - ST]) \leq 4(1 - W), \end{aligned}$$

where $\mathbb{E}_Y[uw] = \mathbb{E}[S(X)T(Z)b(X, Y)c(Z, Y)]$ by independence of the sources and $\mathbb{E}[STbc] \geq \mathbb{E}[abc] - \mathbb{E}[ST - a]$ because $|bc| \leq 1$. Hence $\mathbb{E}_Y[1 - vu] \leq 4(1 - W)$ and $\mathbb{E}_Y[1 - vw] \leq 4(1 - W)$. By (14) with $p = b(X, Y)$ and $r = S(X)v(Y)$,

$$\mathbb{E}[b - Sv] \leq \mathbb{E}[1 - bSv] = \mathbb{E}_Y[1 - vu] \leq 4(1 - W), \quad \text{so} \quad |\mathbb{E}[b] - \mathbb{E}[S] \mathbb{E}[v]| \leq 4(1 - W),$$

and likewise $|\mathbb{E}[c] - \mathbb{E}[T] \mathbb{E}[v]| \leq 4(1 - W)$. Write $\sigma = \mathbb{E}[S]$, $\tau = \mathbb{E}[T]$, $\nu = \mathbb{E}[v]$. If all three means were below $-4(1 - W)$, then $\sigma\tau \leq \mathbb{E}[a] + (1 - W) < 0$, $\sigma\nu < 0$ and $\tau\nu < 0$, whose product $(\sigma\tau\nu)^2$ is nonnegative, a contradiction. $\square$

For $W = 1$ the lemma recovers the consequence of Lemma 3.18 that at least one mean is nonnegative. A compatible law at total-variation distance δ from a parity-perfect law has $1 - W \leq 2\delta$, so the lemma replaces the repair step of [3]: we construct no repaired model, and the bound holds for arbitrary latent alphabets and stochastic responses.

Lemma 3.20 (Distance of the cycle family). *Let $m \geq 3$ and let $q > 0$ be such that $P_{m,q}$ is a probability law. Then $d_{\text{TV}}(P_{m,q}, \mathcal{C}_{C_m}) \geq q/10$; in particular $P_{m,q} \notin \mathcal{C}_{C_m}$.*

Proof. As illustrated in Figure 3, partition $\{0, \dots, m-1\}$ into three nonempty contiguous arcs $\Lambda_1, \Lambda_2, \Lambda_3$ and let $\varphi(o) = (V_1, V_2, V_3)$ with $V_r = \prod_{v \in \Lambda_r} o_v$. If $R \in \mathcal{C}_{C_m}$, then $\varphi_* R \in \mathcal{C}_\Delta$: the three sources on the arc boundaries are independent, each arc product is a function of the two boundary sources adjacent to its arc together with the sources internal to that arc, and we absorb those internal sources into the private randomness of the arc. Each $\partial\Lambda_r$ has two elements, and $\partial(\Lambda_1 \cup \Lambda_2) = \partial\Lambda_3$, so (9) gives

$$\varphi_* P_{m,q} = \Pi(-q, -q, -q).$$

Let $R \in \mathcal{C}_{C_m}$ and $\delta = d_{\text{TV}}(P_{m,q}, R)$. Total variation does not increase under φ , so $Q = \varphi_* R \in \mathcal{C}_\Delta$ satisfies $d_{\text{TV}}(\Pi(-q, -q, -q), Q) \leq \delta$. Since $|\mathbb{E}_P f - \mathbb{E}_Q f| \leq 2d_{\text{TV}}(P, Q)$ for $|f| \leq 1$, the parity $W = \mathbb{E}_Q[V_1 V_2 V_3]$ satisfies $1 - W \leq 2\delta$ and each mean $\mathbb{E}_Q V_r$ is at most $-q + 2\delta$. Lemma 3.19 gives $-q + 2\delta \geq -4(1 - W) \geq -8\delta$, that is $\delta \geq q/10$. $\square$




Figure 3: Reducing cycle incompatibility to triangle parity rigidity. Partition any cycle into three nonempty contiguous arcs and multiply the observations within each arc. The three boundary sources form a triangle; internal sources become private randomness. The target reduces to $\Pi(-q, -q, -q)$, and total variation cannot increase under this map.

Theorem 3.21 (Cycles at every order). *Let $m \geq 3$ and $t \geq 1$, put*

$$q = \frac{1}{4m^2 t^2}, \quad \eta = \frac{q}{32m},$$

and let $\tilde{P}_{m,q} = T_\eta^{\otimes m} P_{m,q}$, where T_η flips a sign with probability η . Then $\tilde{P}_{m,q}$ is a strictly positive rational law with

$$\tilde{P}_{m,q} \in \mathcal{I}_t^{\text{exp}}(C_m) = \mathcal{I}_t^{\text{AI}}(C_m), \quad d_{\text{TV}}(\tilde{P}_{m,q}, \mathcal{C}_{C_m}) \geq \frac{11q}{160} = \frac{11}{640 m^2 t^2} > 0.$$

Consequently no finite order of any of the three hierarchies characterizes $\mathcal{C}_{C_m}$.

Proof. Apply T_η independently to every copied observation of the witness of Lemma 3.17. The resulting law is strictly positive, still invariant under the copy-index permutations, and its characters are those of (13) multiplied by $(1 - 2\eta)^{|F|}$. The same factor multiplies the target characters, and $|F|$ adds over disjoint blocks, so every injectable marginal and every ancestrally

independent product of Lemma 3.17 persists with $\tilde{P}_{m,q}$ in place of $P_{m,q}$, as does the diagonal law. Lemma 3.4 again supplies the recursive prescriptions.

For the distance, $d_{\text{TV}}(\tilde{P}_{m,q}, P_{m,q}) \leq m\eta = q/32$ by a union bound over the flipped coordinates, and Lemma 3.20 gives $d_{\text{TV}}(\tilde{P}_{m,q}, \mathcal{C}_{C_m}) \geq q/10 - q/32 = 11q/160$. As t is arbitrary and $\tilde{P}_{m,q} \notin \mathcal{C}_{C_m}$, the feasible set at order t is strictly larger than $\mathcal{C}_{C_m}$ for every t . $\square$

For each fixed cycle this gives a lower bound of order t^{-2} on the worst-case distance. Section 4 improves the lower bound to order t^{-1} for the triangle and square.

3.5 The five-observer path

For the five-observer path we use a bilocal inequality in place of the cycle parity obstruction. The scenario is

$$A(X), \quad B(X, L), \quad C(L, R), \quad D(R, Z), \quad E(Z),$$

with X, L, R, Z mutually independent (Figure 4). We write $\mathcal{C}_{P_5}$ for its compatible set, record the endpoint observations A, E as bits x, z and the middle observations B, C, D as signs. The order- t inflation of Definition 2.2 has copies $X_a, L_i, R_j, Z_e, a, i, j, e \in [t]$, and the $3t^2 + 2t$ copied observations

$$A^a, \quad B^{ai}, \quad C^{ij}, \quad D^{je}, \quad E^e.$$




Figure 4: The five-observer path. Conditioning on $A = x$ and $E = z$ changes only the outer sources X, Z ; the inner sources L, R remain independent, which gives the bilocal inequality. The target violates it by $\sqrt{1+h} - 1$ and has full support.

3.5.1 The target

Definition 3.22 (Path target). For $0 < h < 1$ define the law P_h by

$$P_h(x, b, c, d, z) = \frac{1}{32} \left[1 + bcd \frac{1+h}{4} (1 + (-1)^{x+z}) \right], \quad x, z \in \{0, 1\}, \quad b, c, d \in \{-1, 1\}. \quad (16)$$

Lemma 3.23 (Properties of P_h). P_h is a probability law, every atom is at least $(1-h)/64$, the pair (A, E) consists of independent fair bits, every proper nonempty moment of (B, C, D) conditional on (A, E) vanishes, and

$$f_{xz} := \mathbb{E}[BCD \mid A = x, E = z] = \frac{1+h}{4} (1 + (-1)^{x+z}) = \begin{cases} (1+h)/2, & x = z, \\ 0, & x \neq z. \end{cases} \quad (17)$$

Proof. The bracket in (16) lies between $1 - (1 + h)/2 = (1 - h)/2$ and $1 + (1 + h)/2$, so every atom lies in $[(1 - h)/64, (3 + h)/64]$ and is positive. Summing bcd over $\{-1, 1\}^3$ gives zero, so the atoms sum to $32 \cdot \frac{1}{32} = 1$ and the (A, E) marginal is uniform on four cells. Given (x, z) , the conditional density of (b, c, d) is $\frac{1}{8}[1 + bcd f_{xz}]$, whose only nonvanishing nontrivial character is bcd , with value f_{xz} ; that is (17) and the vanishing of the proper moments. $\square$

3.5.2 A bilocal obstruction

Lemma 3.24 (Bilocal inequality on the path). *Let $R \in \mathcal{C}_{P_5}$ with $R(A = x, E = z) > 0$ for all four (x, z) , put $f_{xz} = \mathbb{E}_R[BCD \mid A = x, E = z]$ and*

$$I = \frac{1}{4} \sum_{x,z} f_{xz}, \quad J = \frac{1}{4} \sum_{x,z} (-1)^{x+z} f_{xz}.$$

Then $\sqrt{|I|} + \sqrt{|J|} \leq 1$.

Proof. Absorb the response seeds into incident sources. Conditioning on $\{A = x\}$ reweights the law of X alone, and conditioning on $\{E = z\}$ reweights the law of Z alone, so under the conditional law X, L, R, Z are still mutually independent, and L, R keep their original laws. Writing

$$b_x(\ell) = \mathbb{E}[B \mid L = \ell, A = x], \quad d_z(\rho) = \mathbb{E}[D \mid R = \rho, E = z], \\ c(\ell, \rho) = \mathbb{E}[C \mid L = \ell, R = \rho],$$

all with values in $[-1, 1]$, conditional independence of the three responses given the sources gives

$$f_{xz} = \mathbb{E}_{L,R}[b_x(L) c(L, R) d_z(R)].$$

Put $\beta_{\pm} = \frac{1}{2}(b_0 \pm b_1)$ and $\delta_{\pm} = \frac{1}{2}(d_0 \pm d_1)$. Then

$$I = \mathbb{E}_{L,R}[\beta_+(L) c(L, R) \delta_+(R)], \quad J = \mathbb{E}_{L,R}[\beta_-(L) c(L, R) \delta_-(R)],$$

so, using $|c| \leq 1$ and the independence of L and R ,

$$|I| \leq p_+ r_+, \quad |J| \leq p_- r_-, \quad p_{\pm} = \mathbb{E}_L |\beta_{\pm}|, \quad r_{\pm} = \mathbb{E}_R |\delta_{\pm}|.$$

For real u, v with $|u|, |v| \leq 1$, $|\frac{u+v}{2}| + |\frac{u-v}{2}| = \max(|u|, |v|) \leq 1$, so $p_+ + p_- \leq 1$ and $r_+ + r_- \leq 1$. Cauchy–Schwarz gives

$$\sqrt{|I|} + \sqrt{|J|} \leq \sqrt{p_+ r_+} + \sqrt{p_- r_-} \leq \sqrt{p_+ + p_-} \sqrt{r_+ + r_-} \leq 1. \quad \square$$

Fritz [9, Theorem 2.9] identifies classical correlations on P_5 with bilocal conditional distributions after conditioning on the endpoint observations. Inequalities of this shape go back to Branciard, Rosset, Gisin and Pironio [4] for the bilocal scenario with measurement inputs; Lemma 3.24 is the specialization to a single fixed setting, one binary central output and two binary endpoint observations, proved here for arbitrary latent alphabets and stochastic responses.

Corollary 3.25 (Quantitative incompatibility). *For every $0 < h < 1$,*

$$P_h \notin \mathcal{C}_{P_5} \quad \text{and} \quad d_{\text{TV}}(P_h, \mathcal{C}_{P_5}) \geq \frac{h}{96}.$$

Proof. By (17), $I = J = (1 + h)/4$ for P_h , so $\sqrt{I} + \sqrt{J} = \sqrt{1 + h} > 1$ and Lemma 3.24 rules out compatibility.

Let $R \in \mathcal{C}_{P_5}$ and $d = d_{TV}(P_h, R)$. If $d \geq 1/8$ then $d > 1/96 > h/96$, so assume $d < 1/8$. Each endpoint cell satisfies $|R(x, z) - \frac{1}{4}| \leq d$, hence $R(x, z) \geq \frac{1}{4} - d > \frac{1}{8} > 0$ and Lemma 3.24 applies to R . Write $N_R(x, z) = \mathbb{E}_R[BCD \mathbf{1}\{A = x, E = z\}]$ and N_P for the same quantity under P_h , so $f_{xz}^R = N_R(x, z)/R(x, z)$ and likewise for P_h . The eight atoms of one endpoint cell contribute at most $2d$ to $\sum |R - P_h|$, so $|N_R - N_P| \leq 2d$, while $|R(x, z) - P_h(x, z)| \leq d$. Since $|N_P| \leq P_h(x, z) = \frac{1}{4}$,

$$|f_{xz}^R - f_{xz}^P| \leq \frac{|N_R - N_P|}{R(x, z)} + \frac{|N_P| |P_h(x, z) - R(x, z)|}{R(x, z) P_h(x, z)} \leq \frac{2d + d}{R(x, z)} \leq 24d.$$

Averaging, $|I_R - I_P| \leq 24d$ and $|J_R - J_P| \leq 24d$. If $24d < h/4$ then $I_R, J_R > \frac{1+h}{4} - \frac{h}{4} = \frac{1}{4}$, so $\sqrt{|I_R|} + \sqrt{|J_R|} > 1$, contradicting Lemma 3.24. Hence $24d \geq h/4$. $\square$

3.5.3 The witness at every order

Theorem 3.26 (The path at every order). *For every $t \geq 1$ put $\gamma = 1/(4t)$ and $h_t = \gamma^2 = 1/(16t^2)$. Then*

$$P_{h_t} \in \mathcal{I}_t^{\text{exp}}(P_5) = \mathcal{I}_t^{\text{AI}}(P_5), \quad P_{h_t} \notin \mathcal{C}_{P_5},$$

with a strictly positive rational target and a rational witness. Consequently $t_{\min}^H(P_{h_t}) > t$ for $H \in \{\text{NW}, \text{AI}, \text{exp}\}$, no finite order of any of the three hierarchies characterizes $\mathcal{C}_{P_5}$, and

$$d_{TV}(P_{h_t}, \mathcal{C}_{P_5}) \geq \frac{1}{1536 t^2}.$$

Proof. The auxiliary law. On signs $u = (u_1, \dots, u_t)$ and $v = (v_1, \dots, v_t)$ set

$$H_t(u, v) = 2^{-2t} \operatorname{Re} \left[\prod_{i=1}^t (1 + i\gamma u_i) \prod_{j=1}^t (1 - i\gamma v_j) \right]. \quad (18)$$

Writing $\sigma = (u_1, \dots, u_t, -v_1, \dots, -v_t) \in \{-1, 1\}^{2t}$, the bracket is $\prod_{k=1}^{2t} (1 + i\gamma \sigma_k)$, whose real part keeps the even characters with coefficient $(i\gamma)^{|S|} = (-h_t)^{|S|/2}$, so

$$H_t = 2^{-2t} \sum_{\substack{S \subseteq [2t] \\ |S| \text{ even}}} (-h_t)^{|S|/2} \prod_{k \in S} \sigma_k.$$

This is H_{2t, h_t} of Lemma 3.14 read in the coordinates σ , and $(2t)^2 h_t = 1/4$, so H_t is a probability law with

$$H_t(u, v) \geq \frac{2}{3} 2^{-2t} > 0. \quad (19)$$

Its values are rational.

The witness. Draw $(u, v) \sim H_t$; independently draw fair bits X_a and Z_e , fair signs r_i and s_j , and fair signs ν_{ij} , all mutually independent, $a, i, j, e \in [t]$. With $u^0 = 1$ and $u^1 = u$, define every copied observation by

$$A^a = X_a, \quad B^{ai} = r_i u_i^{X_a}, \quad C^{ij} = \begin{cases} r_i s_j, & u_i = v_j, \\ r_i s_j \nu_{ij}, & u_i \neq v_j, \end{cases} \quad D^{je} = s_j v_j^{Z_e}, \quad E^e = Z_e. \quad (20)$$

Each copied observation is a function of the data attached to its copied parents together with its own private sign: r_i and u_i belong to L_i , and s_j and v_j to R_j . Call the resulting law Γ_t . It is a nonnegative rational probability law. The auxiliary density is strictly positive, but the observed table can have zero entries. As in Section 5.1, Γ_t is a law on the copied observations and does not come from an inflated model, since under H_t the families u and v are correlated.

Formal weights. Give each u_i the complex weight $(1 + i\gamma u_i)/2$ and each v_j the weight $(1 - i\gamma v_j)/2$, and keep the actual probabilities of X, Z, r, s, ν . Write $\mathbb{E}_*$ for the resulting formal expectation, a product functional on the $6t + t^2$ auxiliary coordinates, with $\mathbb{E}_* u_i = i\gamma$ and $\mathbb{E}_* v_j = -i\gamma$. Since the weights of the auxiliary coordinates multiply to a complex number whose real part is H_t , and the other coordinates carry real weights,

$$\mathbb{E}_{\Gamma_t}[F] = \operatorname{Re} \mathbb{E}_*[F] \quad \text{for every real-valued } F. \quad (21)$$

Every copied path has law P_h . Fix (a, i, j, e) and condition on $X_a = x, Z_e = z$, which are fair and independent and reproduce the (A, E) marginal of P_h . The masks r_i and s_j are fair signs independent of everything else and each occurs in exactly one of B^{ai}, D^{je} and in C^{ij} , so every product of a proper nonempty subset of $\{B^{ai}, C^{ij}, D^{je}\}$ carries r_i or s_j to an odd power and has formal expectation zero. In the triple product the masks cancel:

$$B^{ai} C^{ij} D^{je} = u_i^x v_j^z \cdot (\mathbf{1}\{u_i = v_j\} + \nu_{ij} \mathbf{1}\{u_i \neq v_j\}),$$

and $\mathbb{E}_* \nu_{ij} = 0$, so $\mathbb{E}_*[B^{ai} C^{ij} D^{je}] = \mathbb{E}_*[u_i^x v_j^z \mathbf{1}\{u_i = v_j\}]$. The formal weight of $(u_i, v_j) = (+1, +1)$ is $\frac{1+i\gamma}{2} \cdot \frac{1-i\gamma}{2} = \frac{1+h}{4}$, and the weight of $(-1, -1)$ is the same real number, so

$$\mathbb{E}_*[u_i^x v_j^z \mathbf{1}\{u_i = v_j\}] = \frac{1+h}{4} (1 + (-1)^{x+z}) = f_{xz},$$

which is real. Every character of the copied path is thus one of: a character of the fair bits X_a, Z_e alone; a character carrying a proper nonempty subset of $\{B^{ai}, C^{ij}, D^{je}\}$, with formal expectation zero; or a character carrying all three, with formal expectation $\frac{1}{4} \sum_{x,z} \chi(x, z) f_{xz}$. These are real and agree with the corresponding characters of P_h in Lemma 3.23, so by (21)

$$\operatorname{Law}_{\Gamma_t}(A^a, B^{ai}, C^{ij}, D^{je}, E^e) = P_h,$$

and $\mathbb{E}_*[G]$ is real for every function G of a copied path.

Injectable marginals. An injectable set of the path contains at most one copied observation over each of A, B, C, D, E and its copy indices agree on shared sources, so it lies in a copied path (a, i, j, e) with consistent indices, after completing unused indices arbitrarily. Its marginal is therefore the corresponding marginal of P_h , and $\mathbb{E}_*[G]$ is real for every function G of it.

Ancestral products and the diagonal law. Let $F_1, \dots, F_r$ be pairwise ancestrally independent injectable sets. Distinct copied sources carry distinct auxiliary coordinates and distinct masks, and distinct copied observations carry distinct private signs, so the F_k depend on disjoint coordinate families and $\mathbb{E}_*[\prod_k G_k] = \prod_k \mathbb{E}_*[G_k]$ for functions G_k of F_k . Each factor is real by the previous paragraph, hence so is the product, and (21) gives

$$\mathbb{E}_{\Gamma_t}\left[\prod_k G_k\right] = \prod_k \mathbb{E}_*[G_k] = \prod_k \mathbb{E}_{\Gamma_t}[G_k].$$

Taking the G_k to be indicators of atoms shows the blocks are independent with the prescribed marginals. The diagonal rows $(A^\ell, B^{\ell\ell}, C^{\ell\ell}, D^{\ell\ell}, E^\ell)$, $\ell \in [t]$, are pairwise ancestrally independent injectable copied paths, so the diagonal law is $P_h^{\otimes t}$.

Symmetry. The density (18) is symmetric in $u_1, \dots, u_t$ and in $v_1, \dots, v_t$ separately, the families (X_a) , (Z_e) , (r_i) , (s_j) and (ν_{ij}) are i.i.d., and (20) is equivariant, so Γ_t is invariant under independent permutations of the four copy-index families. Together with the previous two paragraphs this places P_{h_t} in $\mathcal{I}_t^{\text{AI}}(P_5)$, and Lemma 3.4 gives $\mathcal{I}_t^{\text{exp}}(P_5) = \mathcal{I}_t^{\text{AI}}(P_5)$.

Conclusion. Corollary 3.25 gives $P_{h_t} \notin \mathcal{C}_{P_5}$ and $d_{\text{TV}}(P_{h_t}, \mathcal{C}_{P_5}) \geq h_t/96 = 1/(1536t^2)$. Membership at order t forces $t_{\min}^H(P_{h_t}) > t$ for each of the three hierarchies, by Proposition 2.6. $\square$

Along this family the first rejecting order grows without bound while the target stays uniformly inside the simplex: since $h_t \leq 1/16$, Lemma 3.23 gives

$$P_{h_t}(x, b, c, d, z) \geq \frac{1 - h_t}{64} \geq \frac{15}{1024}$$

for every t and every atom. The theorem gives a distance lower bound of order t^{-2} and leaves the distance asymptotics of this family undetermined. The general upper bound is $O(t^{-1/2})$. It remains open whether the path admits accepted laws at distance $\Omega(t^{-1})$, as the triangle and square do.

4 Quantitative convergence on the triangle and square

We bound the distance from finite-order feasible laws to compatibility on the triangle and square. Modifying the unconstrained Fourier coefficients of the parity witness improves the lower bound from $\Omega(t^{-2})$ to $\Omega(t^{-1})$. A convex-order argument gives an $O(t^{-1/2})$ upper bound. We then compare the hierarchy orders and certify finite threshold gaps.

4.1 A corrected density and a linear lower bound

The density (7) fixes every auxiliary character, including characters that no prescription refers to. Changing those characters restores positivity at a larger q .

Lemma 4.1 (Phase telescoping). *For real $\theta_1, \dots, \theta_m$,*

$$1 - \cos\left(\sum_{g=1}^m \theta_g\right) \leq m \sum_{g=1}^m (1 - \cos \theta_g).$$

Proof. $1 - \cos \phi = \frac{1}{2}|1 - e^{i\phi}|^2$, and $1 - \prod_g e^{i\theta_g} = \sum_g (\prod_{g' < g} e^{i\theta_{g'}})(1 - e^{i\theta_g})$, so $|1 - \prod_g e^{i\theta_g}| \leq \sum_g |1 - e^{i\theta_g}|$. Cauchy–Schwarz gives $(\sum_g a_g)^2 \leq m \sum_g a_g^2$. $\square$

Lemma 4.2 (Corrected density). *Fix $m \in \{3, 4\}$ and $t \geq 1$, and let $q > 0$ satisfy*

$$R := (1 + q)^{t/2} \leq \frac{m^2}{m^2 - 1}, \quad \text{equivalently} \quad q \leq \begin{cases} (9/8)^{2/t} - 1, & m = 3, \\ (16/15)^{2/t} - 1, & m = 4. \end{cases} \quad (22)$$

Set $c = mR^{m-1}$. On m families of t signs $s_{g,i}$ put

$$f_g = \prod_{i=1}^t (1 + i\sqrt{q} s_{g,i}), \quad W(s) = \text{Re} \prod_{g=1}^m f_g + c \sum_{g=1}^m (1 - \text{Re} f_g). \quad (23)$$

Then $W \geq R^{m-1}(m^2 - (m^2 - 1)R) \geq 0$, the law $\rho = 2^{-mt}W$ is a probability law, and its characters are

$$\mathbb{E}_\rho \left[\prod_{w \in S} s_w \right] = \begin{cases} 0, & |S| \text{ odd}, \\ (1-c)(-q)^{|S|/2}, & S \neq \emptyset \text{ even, inside one family}, \\ (-q)^{|S|/2}, & |S| \text{ even, meeting at least two families.} \end{cases} \quad (24)$$

The range (22) contains $0 < q \leq 1/(8t)$ for $m = 4$ and $0 < q \leq 2/(9t)$ for $m = 3$.

Proof. Here $i^2 = -1$. Taking real parts keeps only even powers of $i\sqrt{q}$, so W is real. Since $|1 + i\sqrt{q}s| = (1+q)^{1/2}$ we may write $f_g = Re^{i\theta_g}$, and then

$$\begin{aligned} W &= R^m \cos\left(\sum_g \theta_g\right) + c \sum_g (1 - R \cos \theta_g) \\ &= R^m - cm(R-1) - R^m \left(1 - \cos \sum_g \theta_g\right) + cR \sum_g (1 - \cos \theta_g). \end{aligned}$$

Lemma 4.1 gives, for every $q \geq 0$ and every real c ,

$$W \geq R^m - cm(R-1) + (cR - mR^m) \sum_g (1 - \cos \theta_g).$$

With $c = mR^{m-1}$ the last coefficient vanishes, so $W \geq R^m - m^2 R^{m-1}(R-1) = R^{m-1}(m^2 - (m^2 - 1)R)$, which is nonnegative by (22). Since $(1+x/t)^t \leq e^x$, the range contains $q \leq 2 \ln(16/15)/t$ and $q \leq 2 \ln(9/8)/t$, hence the stated rational ranges.

For the characters, work under the uniform law on the sign cube, where the families are independent and $\mathbb{E}f_g = 1$; the constant c does not depend on s . For S with parts S_g in the families,

$$\mathbb{E} \left[\prod_g f_g \prod_{w \in S} s_w \right] = \prod_g \mathbb{E} \left[f_g \prod_{w \in S_g} s_w \right] = (i\sqrt{q})^{|S|},$$

whose real part is 0 for odd $|S|$ and $(-q)^{|S|/2}$ for even $|S|$. Also $\mathbb{E}[(1 - \operatorname{Re} f_g) \prod_{w \in S} s_w]$ vanishes unless S is a nonempty subset of family g , in which case it equals $-\operatorname{Re}(i\sqrt{q})^{|S|}$. Adding the contributions gives (24); the case $S = \emptyset$ gives $\mathbb{E}_\rho 1 = 1$. $\square$

Theorem 4.3 (The square at $q = \Theta(1/t)$). *For every $t \geq 1$ and every $0 < q \leq (16/15)^{2/t} - 1$, in particular every $0 < q \leq 1/(8t)$, the square parity law P_q of (11) lies in $\mathcal{I}_t^{\text{AI}}(\square) = \mathcal{I}_t^{\text{exp}}(\square)$.*

Proof. Take $m = 4$ in Lemma 4.2, name the four families x, y, z, w after the sources X, Y, Z, W , and define every copied observation by

$$A^{il} = x_i w_l, \quad B^{ij} = x_i y_j, \quad C^{jk} = y_j z_k, \quad D^{kl} = z_k w_l,$$

that is by (12) for $m = 4$. The density (23) is symmetric within each family, so the law is invariant under independent permutations of the copy indices of each source. By Lemma 3.16 every boundary occurring in an AI prescription, and every boundary of a subset of one of its blocks, is empty or meets at least two families; by (24) such characters take the value $(-q)^{|\partial|/2}$, unchanged from (13). The computation of Lemma 3.17 therefore applies verbatim: every injectable set has the marginal of $P_{4,q} = P_q$, ancestrally independent injectable blocks are independent, and the diagonal law is $P_q^{\otimes t}$. The moments that Lemma 4.2 changed are the nonempty even characters inside a single family, and no prescription refers to one. Lemma 3.4 gives the recursive prescriptions. $\square$

Theorem 4.4 (The triangle at $q = \Theta(1/t)$). *For every $t \geq 1$ and every $0 < q \leq (9/8)^{2/t} - 1$, in particular every $0 < q \leq 2/(9t)$, the parity-even law $\Pi(-q, -q, -q)$ of (10) lies in $\mathcal{I}_t^{\text{AI}}(\Delta) = \mathcal{I}_t^{\text{exp}}(\Delta)$.*

Proof. Take $m = 3$ in Lemma 4.2, name the three families x, y, z after the sources X, Y, Z , and define

$$A^{ij} = x_i z_j, \quad B^{ik} = x_i y_k, \quad C^{jk} = z_j y_k,$$

which is (12) for $m = 3$ in the index convention of Section 2.3. Positivity and normalization are Lemma 4.2. Symmetry within each family is immediate. By Lemma 3.16 every character prescribed by an injectable marginal, by an ancestrally independent product, or by the diagonal law has boundary empty or meeting at least two families, so (24) returns $(-q)^{|\partial|/2}$ and the argument of Lemma 3.17 applies unchanged. The original copied triangle has $\mathbb{E}A = \mathbb{E}[x_i z_j] = -q$ and likewise $\mathbb{E}B = \mathbb{E}C = -q$, while $ABC = x_i z_j x_i y_k z_j y_k = 1$, so its law is $\Pi(-q, -q, -q) = P_{3,q}$. Lemma 3.4 gives the recursive prescriptions. $\square$

The exact certificates of Appendix C and the Lean declarations of Section 8 treat the witnesses at $q = 1/(16t)$, which lies in both ranges, with the constants $c = 5$ for the square and $c = 4$ for the triangle. For $tq \leq 1/16$ these constants satisfy $cR \geq mR^m$, so the bound $W \geq R^m - cm(R - 1)$ of the proof holds with them as well.

4.2 Distance and the worst-case profile

For $H \in \{\text{NW}, \text{AI}, \text{exp}\}$ and a pair-source scenario G set

$$H_t^H(G) = \sup_{P \in \mathcal{I}_t^H(G)} d_{\text{TV}}(P, \mathcal{C}_G), \quad (25)$$

the worst-case total-variation distance from compatibility of a law the order- t test accepts, with d_{TV} half the ℓ^1 distance as in Section 2.

Lemma 4.5 (Max-moment inequality for the square). *Every $P \in \mathcal{C}_{\square}$ in sign notation satisfies*

$$\max\{\mathbb{E}A, \mathbb{E}B, \mathbb{E}[AB]\} + 4P(ABCD = -1) \geq 0.$$

Proof. Absorb the response seeds into incident sources, so $A = A(X, W)$, $B = B(X, Y)$, $C = C(Y, Z)$, $D = D(Z, W)$ are deterministic signs and X, Y, Z, W are independent. Fix z and write $g_z(Y) = C(Y, z)$, $h_z(W) = D(z, W)$ and

$$\delta_z = \Pr(A(X, W)B(X, Y) \neq g_z(Y)h_z(W)),$$

so that $\mathbb{E}_Z \delta_Z = P(ABCD = -1) =: \delta$. Put

$$u_z(x) = \mathbb{E}_W[A(x, W)h_z(W)], \quad v_z(x) = \mathbb{E}_Y[B(x, Y)g_z(Y)].$$

Since Y and W are independent, $\mathbb{E}[ABg_zh_z \mid X = x] = u_z(x)v_z(x)$, and as ABg_zh_z is a sign, $\mathbb{E}_X[1 - u_zv_z] = 2\delta_z$.

Let $f_z(x) = \text{sgn}(u_z(x) + v_z(x))$, with $\text{sgn } 0 = 1$, and compare A and B with $A'_z = f_z(X)h_z(W)$ and $B'_z = f_z(X)g_z(Y)$. Then

$$e_z := \Pr(A \neq A'_z) + \Pr(B \neq B'_z) = \mathbb{E}_X\left[1 - \frac{1}{2}|u_z + v_z|\right] \leq \mathbb{E}_X[1 - u_zv_z] = 2\delta_z,$$

using $|u + v| \geq 2uv$ on $[-1, 1]^2$. Independence of X, Y, W gives

$$(\mathbb{E}A'_z)(\mathbb{E}B'_z)(\mathbb{E}[A'_z B'_z]) = (\mathbb{E}f_z)^2(\mathbb{E}g_z)^2(\mathbb{E}h_z)^2 \geq 0,$$

so at least one of the three comparison moments is nonnegative. Each differs from the corresponding moment of P by at most $2e_z$, since $|\mathbb{E}U - \mathbb{E}V| \leq 2\Pr(U \neq V)$ for signs and $\Pr(AB \neq A'_z B'_z) \leq \Pr(A \neq A'_z) + \Pr(B \neq B'_z)$. Hence

$$\max\{\mathbb{E}A, \mathbb{E}B, \mathbb{E}[AB]\} \geq -2e_z \geq -4\delta_z.$$

The left side does not depend on z ; averaging over z finishes the proof. $\square$

Corollary 4.6 (Survivors on the square). $d_{\text{TV}}(P_q, \mathcal{C}_\square) \geq q/6$ for $0 < q \leq 3 - 2\sqrt{2}$. With flip probability $0 < \eta < 1/2$ applied to each of the four bits and to every copied observation, $\tilde{P}_q = T_\eta^{\otimes 4} P_q$ is strictly positive, lies in $\mathcal{I}_t^{\text{AI}}(\square) = \mathcal{I}_t^{\text{exp}}(\square)$ whenever P_q does, and satisfies $d_{\text{TV}}(\tilde{P}_q, \mathcal{C}_\square) \geq q/6 - 4\eta$. At $\underline{q}_t = (16/15)^{2/t} - 1$,

$$H_t^{\text{exp}}(\square) = H_t^{\text{AI}}(\square) \geq \frac{\underline{q}_t}{6} > \frac{1}{47t},$$

and the same bound holds for the supremum over strictly positive accepted laws.

Proof. If $R \in \mathcal{C}_\square$ and $d = d_{\text{TV}}(P_q, R)$, each of $\mathbb{E}_R A, \mathbb{E}_R B, \mathbb{E}_R[AB]$ is at most $-q + 2d$ and the parity error of R is at most d , since P_q has none. Lemma 4.5 gives $0 \leq -q + 2d + 4d$, so $d \geq q/6$. The noise argument of Theorem 3.21 applies to the witness of Theorem 4.3, and $d_{\text{TV}}(\tilde{P}_q, P_q) \leq 4\eta$. Theorem 4.3 places $P_{\underline{q}_t}$ in $\mathcal{I}_t^{\text{AI}}(\square)$, so $H_t^{\text{AI}}(\square) \geq \underline{q}_t/6$, and letting $\eta \downarrow 0$ gives the bound over strictly positive laws. Finally $\underline{q}_t \geq 2 \ln(16/15)/t$ and $\ln(16/15)/3 > 1/47$; Lemma 3.4 identifies the two suprema. $\square$

Corollary 4.7 (Survivors on the triangle). $d_{\text{TV}}(\Pi(-q, -q, -q), \mathcal{C}_\triangle) \geq q/10$ for $0 < q \leq 1/3$. With flip probability $0 < \eta < 1/2$ on each of the three bits and on every copied observation, the noisy target is strictly positive, lies in $\mathcal{I}_t^{\text{AI}}(\triangle) = \mathcal{I}_t^{\text{exp}}(\triangle)$ whenever the target does, and is at distance at least $q/10 - 3\eta$ from $\mathcal{C}_\triangle$. At $\underline{q}_t = (9/8)^{2/t} - 1$,

$$H_t^{\text{exp}}(\triangle) = H_t^{\text{AI}}(\triangle) \geq \frac{\underline{q}_t}{10} > \frac{1}{43t},$$

and the same bound holds for the supremum over strictly positive accepted laws.

Proof. The distance bound is Lemma 3.20 for $m = 3$, whose proof uses only Lemma 3.19 on the triangle itself, and the three flips cost at most 3η in total variation. Feasibility is Theorem 4.4 with the noise argument of Theorem 3.21. At $\underline{q}_t$ this gives $H_t^{\text{AI}}(\triangle) \geq \underline{q}_t/10$, with $\eta \downarrow 0$ for strictly positive laws, and $\underline{q}_t \geq 2 \ln(9/8)/t$ with $\ln(9/8)/5 > 1/43$. $\square$

4.3 A convex-order upper bound

Figure 5 summarizes the two ways of sampling a table used in the proof.

Theorem 4.8 (Convex order along the diagonal). *Let G be a correlation scenario with finitely many independent latent sources and finite-alphabet observations, each a stochastic function of its parent sources, with K joint outcomes. Give each source-free observation an independent private*

source, preserving $\mathcal{C}_G$, and define inflation as in Section 2 with sources in place of edges. For $n \geq 1$, let $P \in \mathcal{I}_n^{\text{NW}}(G)$ with witness Γ . For a deterministic table ω of copied observations let q_ω be the observed law obtained by sampling one copy index uniformly and independently for each source and reading ω at those indices. Then $q_\omega \in \mathcal{C}_G$ for every ω , and for every convex Φ on the observed simplex

$$\mathbb{E}_\Gamma \Phi(q_\omega) \leq \mathbb{E} \Phi(\hat{P}_n), \quad (26)$$

where $\hat{P}_n$ is the empirical law of n independent draws from P . In particular

$$\mathbb{E}_\Gamma \|q_\omega - P\|_2^2 \leq \frac{1 - \|P\|_2^2}{n}, \quad H_n^{\text{NW}}(G) \leq \min\left\{1, \frac{\sqrt{K-1}}{2\sqrt{n}}\right\}. \quad (27)$$

Proof. Fix ω . Sampling an index $I_e \in [n]$ uniformly and independently for each source e and letting each observed variable v output $\omega_v((I_e)_{e \ni v})$, where $e \ni v$ ranges over the sources v reads, is a model of G with independent sources I_e and deterministic responses, so $q_\omega \in \mathcal{C}_G$.

Draw one uniform permutation π_e of $[n]$ for each source, independently of each other and of $\omega \sim \Gamma$, and set

$$T_\ell = \left(\omega_v((\pi_e(\ell))_{e \ni v}) \right)_v, \quad \ell \in [n], \quad \hat{P} = \frac{1}{n} \sum_{\ell=1}^n \delta_{T_\ell}.$$

For each fixed choice of permutations, the symmetry of Γ carries $(T_1, \dots, T_n)$ to the diagonal rows, so $(T_1, \dots, T_n) \sim P^{\otimes n}$; this remains true after averaging over the permutations, so $\hat{P}$ has the law of $\hat{P}_n$. Conditionally on ω and on a fixed ℓ , the indices $(\pi_e(\ell))_e$ are independent and uniform on $[n]$, so $\mathbb{E}[\delta_{T_\ell} \mid \omega] = q_\omega$ and hence $\mathbb{E}[\hat{P} \mid \omega] = q_\omega$. Conditional Jensen gives $\Phi(q_\omega) \leq \mathbb{E}[\Phi(\hat{P}) \mid \omega]$, and averaging over ω gives (26).

Take $\Phi(\mu) = \|\mu - P\|_2^2$. Then $\mathbb{E}\|\hat{P}_n - P\|_2^2 = \frac{1}{n} \sum_x P(x)(1 - P(x)) = (1 - \|P\|_2^2)/n$, which is the first part of (27). Some ω attains at most the mean, and its q_ω is compatible, so

$$d_{\text{TV}}(P, \mathcal{C}_G) \leq \frac{1}{2} \|P - q_\omega\|_1 \leq \frac{\sqrt{K}}{2} \|P - q_\omega\|_2 \leq \frac{\sqrt{K}}{2} \sqrt{\frac{1 - \|P\|_2^2}{n}} \leq \frac{\sqrt{K-1}}{2\sqrt{n}},$$

using $\|P\|_2^2 \geq 1/K$. Total variation never exceeds 1. $\square$




Figure 5: The conditional expectation behind the convex-order bound. The same table yields an empirical distribution through permuted diagonal rows and a compatible law through independent source indices. Conditional Jensen's inequality compares their convex functionals. Squared Euclidean distance then gives the source-count-independent bound.

The conditional-expectation argument removes the source-count factor from the collision-based Euclidean rate $O(\sqrt{L/n})$ for L sources in [15, eqs. (37) and (A7)]. The full diagonal law yields (26) for every convex Φ . For $n \geq 2$, degree-two diagonal constraints already make distinct rows pairwise independent with marginal P , which suffices for the variance calculation and quadratic bound (27).

Corollary 4.9 (Brackets for the square and the triangle). *For every $t \geq 1$,*

$$\frac{1}{6} \left(\left(\frac{16}{15} \right)^{2/t} - 1 \right) \leq H_t^{\text{exp}}(\square) = H_t^{\text{AI}}(\square) \leq H_t^{\text{NW}}(\square) \leq \min \left\{ 1, \frac{\sqrt{15}}{2\sqrt{t}} \right\},$$

$$\frac{1}{10} \left(\left(\frac{9}{8} \right)^{2/t} - 1 \right) \leq H_t^{\text{exp}}(\triangle) = H_t^{\text{AI}}(\triangle) \leq H_t^{\text{NW}}(\triangle) \leq \frac{\sqrt{7}}{2\sqrt{t}}.$$

The two lower bounds exceed $1/(47t)$ and $1/(43t)$.

Proof. Combine Corollaries 4.6 and 4.7 with Theorem 4.8 at $K = 16$ and $K = 8$, and with Proposition 2.6. $\square$

The gap between t^{-1} and $t^{-1/2}$ remains open for both scenarios (Question 7.1).

4.4 Order conversion on the square

Theorem 4.10 (NW and AI differ by a bounded change of order). *For every $t \geq 2$,*

$$\mathcal{I}_{\lfloor 3t/2 \rfloor}^{\text{NW}}(\square) \subseteq \mathcal{I}_t^{\text{AI}}(\square) = \mathcal{I}_t^{\text{exp}}(\square) \subseteq \mathcal{I}_t^{\text{NW}}(\square).$$

The proof in Appendix B.1 places the components of an AI set on at most $\lfloor 3t/2 \rfloor$ diagonal rows.

Hence $H_{\lfloor 3t/2 \rfloor}^{\text{NW}}(\square) \leq H_t^{\text{AI}}(\square) \leq H_t^{\text{NW}}(\square)$, so the two hierarchies have the same polynomial decay exponent if either has one.

4.5 Finite-order separation

For the square parity family P_q and the triangle family $\Pi(-q, -q, -q)$, the accepted parameters at order t form an interval $[0, q_t^H]$ (Lemma B.4). Rational certificates bracket the endpoints through order ten and separate the tests at the following orders.

Proposition 4.11 (The parity thresholds at orders up to ten). *For the square and the triangle, the thresholds $q_t^{\text{AI}} = q_t^{\text{exp}}$ and q_t^{NW} satisfy Table 1 for $1 \leq t \leq 10$. In particular:*

- (i) $q_t^{\text{AI}} < q_t^{\text{NW}}$ at every even $t \leq 10$, so $\mathcal{I}_t^{\text{AI}} \subsetneq \mathcal{I}_t^{\text{NW}}$ at those orders on both scenarios;
 - (ii) the certified brackets agree at every odd $t \leq 9$; agreement of brackets does not establish equality of thresholds;
 - (iii) $q_{t+1}^{\text{NW}} < q_t^{\text{AI}}$ for $2 \leq t \leq 9$, $t \neq 3$, on both scenarios, whereas the brackets at orders three (AI) and four (NW) agree;
 - (iv) on the square, $P_{3/20} \in \mathcal{I}_2^{\text{NW}} \setminus \mathcal{I}_2^{\text{AI}}$ and $P_{1/10} \in \mathcal{I}_2^{\text{AI}} \setminus \mathcal{I}_3^{\text{NW}}$.
- Consequently $\mathcal{I}_3^{\text{NW}}(\square) \subsetneq \mathcal{I}_2^{\text{AI}}(\square) = \mathcal{I}_2^{\text{exp}}(\square) \subsetneq \mathcal{I}_2^{\text{NW}}(\square)$.

The proof is in Appendix B. Figure 6 plots tq_t^H .

The degree sets also give an order conversion along the parity direction on both scenarios.

Proposition 4.12 (Order conversion along the parity direction). *Let $s \geq 1$ and $t \geq \max\{s, \lfloor 3s/2 \rfloor\}$. Then $q_t^{\text{NW}} \leq q_s^{\text{AI}}$ on the square and on the triangle.*

The proof is in Appendix B.4. At $s = 3$ it gives $q_4^{\text{NW}} \leq q_3^{\text{AI}}$, and at these two orders the certified brackets agree, the exception in Proposition 4.11(iii).

At order ten, Table 1 gives $tq_t^H \approx 0.385$ (AI) and 0.401 (NW) on the square, and 0.833 (AI) and 0.870 (NW) on the triangle. These exceed the asymptotic constants $2\ln(16/15) \approx 0.129$ and $2\ln(9/8) \approx 0.236$ in our all-order feasibility guarantees. The finite brackets do not determine whether tq_t^H converges. Exact certificates extend the odd and even pattern of Proposition 4.11 on the triangle: q_t^{AI} and q_t^{NW} have identical certified brackets at $t = 11$ and $t = 13$, with an ancestral-independence dual supported on Navascués–Wolfe keys only, and $q_{12}^{\text{AI}} < q_{12}^{\text{NW}}$. As in Proposition 4.11(ii), identical brackets do not establish equal thresholds. The odd/even pattern remains open beyond the certified orders.




Figure 6: Certified parity thresholds through order ten, rescaled by t . Markers show lower endpoints; two-sided brackets have width $4 \cdot 10^{-9}$ before rescaling. The square NW order-two point is a lower bound only. Dashed: the all-order guarantees $q = (16/15)^{2/t} - 1$ and $q = (9/8)^{2/t} - 1$ of Theorems 4.3 and 4.4. Panels use different vertical scales; connecting segments are visual aids only.

5 The defect family and rejecting-order growth

The parity laws of Sections 3.4 and 4 give nontermination and worst-case lower bounds. This section studies a second triangle family, whose passing and rejecting bounds have the same exponent and whose distance to compatibility we compute to leading order.

For $0 < \varepsilon < 1$ and $0 \leq r \leq 1$ define the three-bit law

$$Q(\varepsilon, r) = \varepsilon \delta_{000} + (1 - \varepsilon) \text{Bern}(r)^{\otimes 3}, \quad (28)$$

a point mass at 000 with weight ε mixed with three independent bits that equal 1 with probability r . Throughout, $\text{Bern}(r)$ denotes the law on $\{0, 1\}$ with $\text{Bern}(r)(1) = r$.

Theorem 5.1 (Membership at every finite order). *Let $t \geq 1$, $0 < \varepsilon < 1$ and $0 \leq r \leq (1 - \varepsilon)^{t-1}$. Then*

$$Q(\varepsilon, r) \in \mathcal{I}_t^{\text{exp}} \subseteq \mathcal{I}_t^{\text{AI}} \subseteq \mathcal{I}_t^{\text{NW}}.$$

Theorem 5.2 (No finite characterizing order). *For every integer $t \geq 1$, with $\varepsilon_t = 1/(2t^3)$ and $r_t = (1 - \varepsilon_t)^{t-1}$, the law $P_t = Q(\varepsilon_t, r_t)$ satisfies*

$$P_t \in \mathcal{I}_t^{\text{exp}} \subseteq \mathcal{I}_t^{\text{AI}} \subseteq \mathcal{I}_t^{\text{NW}} \quad \text{and} \quad P_t(000)^2 - P_{t,A}(0)P_{t,B}(0)P_{t,C}(0) \geq \frac{\varepsilon_t^2}{2} > 0,$$

so $P_t \notin \mathcal{C}_\Delta$. Hence none of the three finite-order tests characterizes $\mathcal{C}_\Delta$ at any finite order.

Theorem 5.2 includes $t = 1$, where $P_1 = \frac{1}{2}(\delta_{000} + \delta_{111})$ is the perfectly correlated three-bit law, the standard example of a triangle-incompatible distribution, with Finner margin $1/8 = \varepsilon_1^2/2$; it passes the trivial order-one test.

Corollary 5.3 (Larger alphabets). *For the triangle with observed alphabets of any sizes $n_A, n_B, n_C \geq 2$, no finite order of the Navascués–Wolfe hierarchy characterizes the compatible set.*

Proof. The binary-support transfer in Remark 3.8 applies to P_t and its inflation witness. $\square$

5.1 The defect cube

Fix $t \geq 1$, $0 < \varepsilon < 1$ and $0 \leq r \leq (1 - \varepsilon)^{t-1}$. The witness is a joint law Γ_t on the $3t^2$ copied observations built from independent bits indexed by the cells of the cube $[t]^3$.

Roots. Let $D_{ijk} \sim \text{Bern}(\varepsilon)$, $(i, j, k) \in [t]^3$, be mutually independent *defect* bits, and let $N_v \sim \text{Bern}(s)$, one for each copied observation v , be mutually independent *private* bits, independent of the defects, with

$$s = 1 - \frac{r}{(1 - \varepsilon)^{t-1}} \in [0, 1]. \quad (29)$$

The endpoint $r = 0$ uses $s = 1$; the endpoint $r = (1 - \varepsilon)^{t-1}$ uses $s = 0$, in which case the private bits play no role.

Outputs. Define

$$\begin{aligned} A^{ij} &= \mathbf{1}\{N_{A^{ij}} = 0 \text{ and } D_{ijk} = 0 \ \forall k\}, \\ B^{ik} &= \mathbf{1}\{N_{B^{ik}} = 0 \text{ and } D_{ijk} = 0 \ \forall j\}, \\ C^{jk} &= \mathbf{1}\{N_{C^{jk}} = 0 \text{ and } D_{ijk} = 0 \ \forall i\}. \end{aligned} \quad (30)$$

Figure 7 shows the line geometry: A^{ij} reads the line of cells $\Lambda(A^{ij}) = \{(i, j, k) : k \in [t]\}$, B^{ik} the line $\Lambda(B^{ik}) = \{(i, j, k) : j \in [t]\}$, and C^{jk} the line $\Lambda(C^{jk}) = \{(i, j, k) : i \in [t]\}$; each output equals 1 exactly when its private bit and every defect on its line are 0. Let Γ_t be the joint law of the outputs. For $s = 0$ its entries are

$$\Gamma_t(w) = \sum_{d \in \{0,1\}^{[t]^3} : F_t(d)=w} \varepsilon^{|d|} (1 - \varepsilon)^{t^3 - |d|}. \quad (31)$$

For rational ε , these entries are rational. Here F_t is the map (30) from defect configurations to observed assignments. The entries are nonnegative and sum to $(\varepsilon + 1 - \varepsilon)^{t^3} = 1$.

Γ_t is a probability law on the copied observations, which is what a finite inflation test asks for; we do not obtain it by running an inflated triangle model.




Figure 7: The defect-cube construction, shown for four indices per axis. A copied observation outputs one precisely when all defects on its line and its private bit are zero. The shared cell supplies the $\varepsilon\delta_{000}$ component; disjoint residual lines supply the product component. The defects give an inflation witness; they are not independent pair sources for the target.

Lemma 5.4 (Disjoint ancestry gives disjoint inputs). *If two sets of copied observations have disjoint copied latent ancestry in the order- t inflation, then their defect lines and their private bits are disjoint, and the two sets of outputs are independent under Γ_t .*

Proof. Two lines of the cube intersect only if they share the index that corresponds to a common source type: $\Lambda(A^{ij}) \cap \Lambda(B^{i'k}) \neq \emptyset$ forces $i = i'$, and then X_i is a common ancestor; similarly $\Lambda(A^{ij}) \cap \Lambda(C^{j'k}) \neq \emptyset$ forces $j = j'$ and $\Lambda(B^{ik}) \cap \Lambda(C^{jk'}) \neq \emptyset$ forces $k = k'$; two lines of the same type are disjoint unless the observations coincide. Taking unions, disjoint ancestry implies disjoint defect supports. Distinct observations have distinct private bits. Outputs that are functions of disjoint families of independent bits are independent. $\square$

Lemma 5.5 (Law of a copied triangle). *For every $(i, j, k) \in [t]^3$, $\text{Law}_{\Gamma_t}(A^{ij}, B^{ik}, C^{jk}) = Q(\varepsilon, r)$.*

Proof. The three lines of Δ_{ijk} meet pairwise exactly at the cell (i, j, k) . Condition on D_{ijk} . If $D_{ijk} = 1$, all three outputs are 0. If $D_{ijk} = 0$, each output depends on its own private bit and on the $t - 1$ remaining cells of its line; these three families are pairwise disjoint, so the outputs are conditionally independent, each equal to 1 with probability $(1 - s)(1 - \varepsilon)^{t-1} = r$ by (29). Averaging over D_{ijk} gives (28). $\square$

Lemma 5.6 (Symmetry). Γ_t is invariant under independent permutations of the X -, Z - and Y -copy indices.

Proof. A permutation σ of the X -indices sends $D_{ijk} \mapsto D_{\sigma(i)jk}$, $N_{A^{ij}} \mapsto N_{A^{\sigma(i)j}}$, and so on; it permutes the two families of i.i.d. bits and carries (30) to itself. The same holds for the other two index families. $\square$

5.2 Incompatibility by Finner

Lemma 5.7 (Finner inequality for the triangle). *Every $P \in \mathcal{C}_{\Delta}$ satisfies $P(000)^2 \leq P_A(0)P_B(0)P_C(0)$.*

Proof. Let the sources have laws μ_X, μ_Y, μ_Z on arbitrary measurable spaces and let

$$f(x, z) = \Pr(A = 0 \mid x, z), \quad g(x, y) = \Pr(B = 0 \mid x, y), \quad h(z, y) = \Pr(C = 0 \mid z, y)$$

be the response probabilities, measurable functions with values in $[0, 1]$. Conditional independence of the responses given the sources gives $P(000) = \int f(x, z)g(x, y)h(z, y) d\mu_X d\mu_Y d\mu_Z$. Put $u(x)^2 = \int g(x, y)^2 d\mu_Y(y)$ and $v(z)^2 = \int h(z, y)^2 d\mu_Y(y)$. Cauchy–Schwarz in y , then Cauchy–Schwarz in the independent pair (x, z) , then $f^2 \leq f$, $g^2 \leq g$, $h^2 \leq h$ (all three functions lie in $[0, 1]$) give

$$\begin{aligned} P(000) &\leq \int f(x, z)u(x)v(z) d\mu_X d\mu_Z \\ &\leq \left(\int f^2 d\mu_X d\mu_Z \right)^{1/2} \left(\int u^2 d\mu_X \int v^2 d\mu_Z \right)^{1/2} \leq \sqrt{P_A(0)P_B(0)P_C(0)}. \end{aligned}$$

Tonelli’s theorem justifies every interchange, since all integrands are nonnegative and bounded. The argument uses neither finiteness of the latent alphabets nor determinism of the responses. $\square$

The inequality is the triangle case of Finner’s generalized Hölder inequality [6]; Renou et al. [19] derived it for the triangle network in this event form and proved it for quantum sources and for no-signalling boxes combined by local wirings.

Lemma 5.8 (Violation by the witness). *Let $t \geq 1$, $0 < \varepsilon < t^{-3}$ and $P = Q(\varepsilon, (1 - \varepsilon)^{t-1})$. Then $P(000)^2 > P_A(0)P_B(0)P_C(0)$; in particular $P \notin \mathcal{C}_\Delta$. For $\varepsilon = \varepsilon_t = 1/(2t^3)$,*

$$P_t(000)^2 - P_{t,A}(0)P_{t,B}(0)P_{t,C}(0) \geq \varepsilon_t^2 - t^3\varepsilon_t^3 = \frac{\varepsilon_t^2}{2}.$$

Proof. With $r = (1 - \varepsilon)^{t-1}$, $P(000) = \varepsilon + (1 - \varepsilon)(1 - r)^3 \geq \varepsilon$ and, by symmetry of the three coordinates, $P_A(0) = P_B(0) = P_C(0) = 1 - (1 - \varepsilon)r = 1 - (1 - \varepsilon)^t \leq t\varepsilon$ by Bernoulli’s inequality. Hence $P(000)^2 \geq \varepsilon^2$ and $P_A(0)P_B(0)P_C(0) \leq t^3\varepsilon^3$, and $\varepsilon^2 > t^3\varepsilon^3$ exactly when $\varepsilon < t^{-3}$. For $\varepsilon_t = 1/(2t^3)$ the difference is at least $\varepsilon_t^2(1 - t^3\varepsilon_t) = \varepsilon_t^2/2$. $\square$

5.3 Survival of every finite order

The defect law Γ_t witnesses Theorem 5.1.

5.3.1 The Navascués–Wolfe conditions

Symmetry is Lemma 5.6. For the diagonal law, let R_ℓ be the union of the three lines of the diagonal triangle $\Delta_{\ell\ell\ell}$:

$$R_\ell = \{(\ell, \ell, k) : k \in [t]\} \cup \{(\ell, j, \ell) : j \in [t]\} \cup \{(i, \ell, \ell) : i \in [t]\}. \quad (32)$$

Equivalently, R_ℓ is the set of cells with at least two coordinates equal to ℓ ; it has $3t - 2$ cells.

Lemma 5.9 (Diagonal supports are mutually disjoint). *If $\ell \neq m$ then $R_\ell \cap R_m = \emptyset$. Consequently the t diagonal triangles depend on mutually disjoint families of independent bits, are mutually independent under Γ_t , and their joint law is $Q(\varepsilon, r)^{\otimes t}$.*

Proof. A cell has three coordinates; it cannot have two of them equal to ℓ and two equal to $m \neq \ell$. The private bits of distinct diagonal triangles are distinct. Mutual independence of functions of disjoint families of independent variables follows, and Lemma 5.5 identifies each factor. $\square$

With symmetry and the tensor-power diagonal law, $Q(\varepsilon, r) \in \mathcal{I}_t^{\text{NW}}$.

5.3.2 The ancestral-independence conditions

Every injectable set is a subset of some copied triangle Δ_{ijk} (Definition 2.4), so its prescribed marginal is a projection of $Q(\varepsilon, r)$, and Lemma 5.5 matches it. By Lemma 5.4 applied to the union, the outputs in a union of pairwise ancestrally independent injectable sets depend on pairwise disjoint families of independent bits, so they are mutually independent with the prescribed product law. This proves $Q(\varepsilon, r) \in \mathcal{I}_t^{\text{AI}}$.

5.3.3 The recursively expressible conditions

Lemma 5.10 (Recursive prescriptions hold). *Every marginal prescribed by Definition 2.5 equals the corresponding marginal of Γ_t .*

Proof. The preceding argument proves the AI prescriptions. The root-sink lemma (Lemma 3.4), applied to the triangle, gives every recursively expressible prescription. $\square$

This proves $Q(\varepsilon, r) \in \mathcal{I}_t^{\text{exp}}$ and completes the proof of Theorem 5.1. Theorem 5.2 follows by combining Theorem 5.1 at $r = r_t = (1 - \varepsilon_t)^{t-1}$ with Lemma 5.8.

5.4 Finite-order fan inequalities

For every law feasible at order t , we bound the probability of the all-zero atom by the one-variable zero marginals.

Theorem 5.11 (Fan inequalities). *Let P be a law on three bits with $a = P_A(0)$, $b = P_B(0)$, $c = P_C(0)$, $z = P(000)$. If $P \in \mathcal{I}_t^{\text{NW}}$ then*

$$t z \leq a + \binom{t}{2} bc \quad \text{and} \quad t(z^2 - abc) \leq az - abc. \quad (33)$$

Consequently, if $z^2 > abc$ (a Finner violation), then

$$t_{\min}^{\text{NW}}(P) \leq \left\lfloor \frac{z \min\{a, b, c\} - abc}{z^2 - abc} \right\rfloor + 1, \quad (34)$$

and the same bound holds for $t_{\min}^{\text{AI}}$ and $t_{\min}^{\text{exp}}$.

Relabelling the local outcomes gives the corresponding bound for a Finner violation at any atom.

Proposition 5.12 (No uniformly divergent distance lower bound). *For $0 < p < 1$ let $R_p = (1 - p)\delta_{111} + p\delta_{000}$. Then $R_p \notin \mathcal{C}_\Delta$, $R_p \rightarrow \delta_{111}$ as $p \downarrow 0$, and $t_{\min}^H(R_p) = 2$ for all three hierarchies and all p . Hence no inequality $t_{\min}^H(P) \geq c d_{\text{TV}}(P, \mathcal{C}_\Delta)^{-\alpha}$ with $c, \alpha > 0$ holds for all incompatible P ; the growth in Proposition 5.13 is a property of the direction of approach to the boundary.*

Proof of Theorem 5.11. For $t = 1$, the two inequalities reduce to $z \leq a$ and $z^2 \leq az$, which hold since $0 \leq z \leq a$. Now let Γ be an order- t feasible law for P , $t \geq 2$, and consider the $2t + 1$ copied observations A^{11}, B^{1k}, C^{1k} , $k \in [t]$: the fan (Figure 8) of copied triangles $\Delta_{11k} = \{A^{11}, B^{1k}, C^{1k}\}$ sharing the vertex A^{11} . Define the indicators $A = \mathbf{1}[A^{11} = 0]$, $B_k = \mathbf{1}[B^{1k} = 0]$, $C_k = \mathbf{1}[C^{1k} = 0]$ and $T_k = AB_k C_k$. Three identities follow from Definition 2.3: $\mathbb{E}A = a$ and $\mathbb{E}T_k = z$, because a




Figure 8: A fan with four copied triangles sharing A^{11} . Solid edges indicate shared ancestry within each triangle; the drawing is an incidence schematic. The cross pair $B^{1k}, C^{1\ell}$ has disjoint ancestors for $k \neq \ell$ and can be placed in distinct diagonal rows. These cross moments yield both inequalities in Theorem 5.11.

permutation of the Y -indices carries Δ_{11k} to the diagonal triangle Δ_{111} , whose law is P ; and for $k \neq \ell$,

$$\mathbb{E}[B_k C_\ell] = bc, \quad (35)$$

because the independent permutations $X_1 \mapsto 1$, $Z_1 \mapsto 2$, $Y_k \mapsto 1$, $Y_\ell \mapsto 2$ (extended arbitrarily to $[t]$) carry $B^{1k}, C^{1\ell}$ to B^{11}, C^{22} , which lie in the two diagonal triangles $\Delta_{111}, \Delta_{222}$ whose joint law is $P \otimes P$ by the degree-two diagonal condition. The argument uses no product prescription beyond the diagonal law, and (35) is not available for $k = \ell$.

First inequality. For every deterministic assignment of the fan,

$$\sum_k A B_k C_k \leq A + \frac{1}{2} \sum_{k \neq \ell} B_k C_\ell. \quad (36)$$

If $A = 0$ the left side vanishes. If $A = 1$, put $S = \sum_k B_k C_k$; then $\sum_k B_k \geq S$ and $\sum_k C_k \geq S$, so the right side is $1 + \frac{1}{2}[(\sum_k B_k)(\sum_k C_k) - S] \geq 1 + \frac{1}{2}(S^2 - S)$, and the right side minus the left side is at least $\frac{1}{2}(S - 1)(S - 2) \geq 0$ for every integer $S \geq 0$. Taking expectations under Γ and using the three identities gives $tz \leq a + \binom{t}{2}bc$.

Second inequality. Let $U = \sum_k T_k$. Since $U = AU$, Cauchy–Schwarz gives $(\mathbb{E}U)^2 \leq \mathbb{E}A \cdot \mathbb{E}U^2$, and $U^2 = \sum_k T_k + \sum_{k \neq \ell} T_k T_\ell$ with $T_k T_\ell \leq B_k C_\ell$. Hence $(tz)^2 \leq a[tz + t(t-1)bc]$, which rearranges to $t(z^2 - abc) \leq az - abc$.

Rejecting order. If $z^2 > abc$, the second inequality fails as soon as $t > (az - abc)/(z^2 - abc)$; we root the fan at whichever of A, B, C has the smallest zero marginal, giving (34). The ratio is at least 1 since numerator minus denominator equals $z(\min\{a, b, c\} - z) \geq 0$, so the bound is never below 2, consistent with order one never rejecting. The smaller feasible sets of (1) reject no later. $\square$

Inequality (36) is a pointwise marginal-polytope certificate in the sense of [22], on a $1 \times 1 \times t$ box of source copies; here we have it in closed form at every order.

Proof of Proposition 5.12. R_p has $a = b = c = z = p$, so $z^2 = p^2 > p^3 = abc$ and Lemma 5.7 gives incompatibility. Order one is passed by every law. At order two the first fan inequality reads $2p \leq p + p^2$, which fails for $p < 1$. Hence $t_{\min}^H(R_p) = 2$ for all three hierarchies, while $d_{\text{TV}}(R_p, \mathcal{C}_\Delta) \leq p \rightarrow 0$. $\square$

5.5 Rejecting-order exponent along the family

Proposition 5.13 (One family with divergent rejecting order). *For $0 < \varepsilon < 1/8$ let $P_\varepsilon = Q(\varepsilon, 1 - \frac{1}{2}\varepsilon^{2/3})$, and write $\sigma = \frac{1}{2}\varepsilon^{2/3}$, $m = \varepsilon + (1 - \varepsilon)\sigma$ for the common one-variable zero marginal and $z = \varepsilon + (1 - \varepsilon)\sigma^3$ for $P_\varepsilon(000)$. Then $P_\varepsilon \notin \mathcal{C}_\Delta$, $P_\varepsilon \rightarrow \delta_{111} \in \mathcal{C}_\Delta$ as $\varepsilon \downarrow 0$, and for $H \in \{\text{exp, AI, NW}\}$ the first rejecting order $t_{\min}^H(P_\varepsilon) = \min\{t : P_\varepsilon \notin \mathcal{I}_t^H\}$ is finite and satisfies*

$$\left\lfloor 1 + \frac{1}{2}\varepsilon^{-1/3} \right\rfloor + 1 \leq t_{\min}^H(P_\varepsilon) \leq \lfloor \tau_-(\varepsilon) \rfloor + 1 \quad (\varepsilon \leq 1/64),$$

$$\tau_- = \frac{z + m^2/2 - \sqrt{(z + m^2/2)^2 - 2m^3}}{m^2},$$

so that

$$\frac{1}{2} \leq \liminf_{\varepsilon \downarrow 0} \varepsilon^{1/3} t_{\min}^H(P_\varepsilon) \leq \limsup_{\varepsilon \downarrow 0} \varepsilon^{1/3} t_{\min}^H(P_\varepsilon) \leq 4 - 2\sqrt{3} \approx 0.536.$$

In particular $t_{\min}^H(P_\varepsilon) = \Theta(\varepsilon^{-1/3})$ for each of the three hierarchies.

Proof of Proposition 5.13. Put $\sigma = \frac{1}{2}\varepsilon^{2/3}$, so $P_\varepsilon = Q(\varepsilon, 1 - \sigma)$, $m = \varepsilon + (1 - \varepsilon)\sigma$ and $z = \varepsilon + (1 - \varepsilon)\sigma^3$. *Incompatibility:* $z \geq \varepsilon$ and $m \leq \varepsilon + \sigma = \varepsilon^{2/3}(\varepsilon^{1/3} + \frac{1}{2})$, so $m^3 < \varepsilon^2 \leq z^2$ because $\varepsilon^{1/3} < \frac{1}{2}$ when $\varepsilon < 1/8$; Lemma 5.7 applies. *Limit:* $\varepsilon \rightarrow 0$ and $1 - \sigma \rightarrow 1$, so $P_\varepsilon \rightarrow \delta_{111}$, which is compatible. *Lower bound:* if $t \leq 1 + \frac{1}{2}\varepsilon^{-1/3}$ then $(t - 1)\varepsilon \leq \sigma$, and Bernoulli's inequality gives $(1 - \varepsilon)^{t-1} \geq 1 - (t - 1)\varepsilon \geq 1 - \sigma$, so Theorem 5.1 applies with $r = 1 - \sigma$ and $P_\varepsilon \in \mathcal{I}_t^{\text{exp}}$; hence every rejecting order exceeds $\lfloor 1 + \frac{1}{2}\varepsilon^{-1/3} \rfloor$. Finiteness is asymptotic completeness together with (1).

Upper bound. Write $v_t = tz - m - \binom{t}{2}m^2$; by the first fan inequality with $a = b = c = m$, positivity of v_t rejects at order t . For $t = x\varepsilon^{-1/3} + O(1)$, $tz = x\varepsilon^{2/3} + O(\varepsilon)$, $m = \frac{1}{2}\varepsilon^{2/3} + O(\varepsilon)$ and $\binom{t}{2}m^2 = \frac{x^2}{8}\varepsilon^{2/3} + O(\varepsilon)$, so

$$v_t = \left(x - \frac{1}{2} - \frac{x^2}{8}\right)\varepsilon^{2/3} + O(\varepsilon),$$

whose leading coefficient is positive exactly for $4 - 2\sqrt{3} < x < 4 + 2\sqrt{3}$. Fixing x slightly above $4 - 2\sqrt{3}$ gives rejection for all small ε , hence the lim sup bound. For the exact finite statement, v_t is the quadratic $-\frac{m^2}{2}t^2 + (z + \frac{m^2}{2})t - m$ in t , with discriminant $D = (z + m^2/2)^2 - 2m^3$ and roots $\tau_\pm = (z + m^2/2 \pm \sqrt{D})/m^2$. For $\varepsilon \leq 1/64$ we have $m \leq \frac{3}{4}\varepsilon^{2/3}$ and $z \geq \varepsilon$, so $D \geq \varepsilon^2 - 2m^3 \geq \varepsilon^2(1 - \frac{27}{32}) = \frac{5}{32}\varepsilon^2$ and $\tau_+ - \tau_- = 2\sqrt{D}/m^2 \geq \frac{4\sqrt{10}}{9}\varepsilon^{-1/3} > 1$; since $v_1 = z - m < 0$ the integer $\lfloor \tau_- \rfloor + 1$ lies strictly between the roots, where $v_t > 0$. Hence $t_{\min}^{\text{NW}}(P_\varepsilon) \leq \lfloor \tau_- \rfloor + 1 = (4 - 2\sqrt{3})\varepsilon^{-1/3} + O(1)$. □

5.6 Distance to the compatible set

We compute the leading distance from P_ε to $\mathcal{C}_\Delta$, which expresses the rejecting-order exponent in terms of distance.

Theorem 5.14 (Distance asymptotics). *For $P_\varepsilon = Q(\varepsilon, 1 - \frac{1}{2}\varepsilon^{2/3})$, $0 < \varepsilon < 1/8$,*

$$d_{\text{TV}}(P_\varepsilon, \mathcal{C}_\Delta) = (1 - 2^{-3/2})\varepsilon + O(\varepsilon^{4/3}), \quad d_2(P_\varepsilon, \mathcal{C}_\Delta) = \sqrt{8/7}(1 - 2^{-3/2})\varepsilon + O(\varepsilon^{4/3}).$$

Corollary 5.15 (Direction-specific boundary complexity). *For $H \in \{\text{NW}, \text{AI}, \text{exp}\}$,*

$$t_{\min}^H(P_\varepsilon) = \Theta(d_{\text{TV}}(P_\varepsilon, \mathcal{C}_\Delta)^{-1/3}) = \Theta(d_2(P_\varepsilon, \mathcal{C}_\Delta)^{-1/3}) \quad (\varepsilon \downarrow 0).$$

This is a statement about one direction of approach to the boundary; by Proposition 5.12 no such lower bound holds uniformly over all incompatible laws.

Proof of Theorem 5.14. Put $\sigma = \frac{1}{2}\varepsilon^{2/3}$, $m = \varepsilon + (1 - \varepsilon)\sigma$, $z = \varepsilon + (1 - \varepsilon)\sigma^3$, $c_0 = 2^{-3/2}$ and $d_0 = 1 - c_0$.

Lower bounds. The product law $\text{Bern}(1 - \sigma)^{\otimes 3}$ is compatible and within total variation ε of P_ε . Compactness of $\mathcal{C}_\Delta$ gives a nearest law in either norm [20]. A total-variation minimizer Q has distance $\delta = O(\varepsilon)$. Since probabilities of events differ by at most δ ,

$$Q(000) \geq z - \delta, \quad Q_V(0) \leq m + \delta \quad (V = A, B, C).$$

Finner gives

$$z - \delta \leq (m + \delta)^{3/2} = c_0\varepsilon + O(\varepsilon^{4/3}),$$

so $\delta \geq d_0\varepsilon - O(\varepsilon^{4/3})$.

A Euclidean minimizer Q_2 also satisfies $\|P_\varepsilon - Q_2\|_2 = O(\varepsilon)$, by comparison with the product law. Cauchy–Schwarz bounds each one-variable marginal error by $2\|P_\varepsilon - Q_2\|_2$, since each marginal sums four atoms. Hence Finner gives $Q_2(000) \leq c_0\varepsilon + O(\varepsilon^{4/3})$. Set $D = P_\varepsilon(000) - Q_2(000)$; then $D \geq d_0\varepsilon - O(\varepsilon^{4/3}) > 0$ for small ε . The other seven coordinates of $Q_2 - P_\varepsilon$ sum to D , so

$$\|Q_2 - P_\varepsilon\|_2^2 \geq D^2 + \frac{D^2}{7}.$$

This proves the Euclidean lower bound.

A compatible approximation. Set $h = d_0/7$ and $u = \sigma + (c_0 + h)\varepsilon$. Each source carries a base bit with success probability $\sqrt{u}$ and a flag with success probability $h\varepsilon$, all six bits independent. A node outputs 0 if both incident base bits equal 1 or either incident flag equals 1. These probabilities lie in $[0, 1]$ for small ε , and the resulting law Q_ε belongs to $\mathcal{C}_\Delta$.

Without flags, the atom probabilities are $u^{3/2}$ at 000, $u - u^{3/2}$ at each one-zero atom, zero at each two-zero atom, and $1 - 3u + 2u^{3/2}$ at 111. A single flag forces the two nodes incident to its source to zero. Except on a base-output event of probability $O(u)$, it transfers mass from 111 to the two-zero atom associated with that source. Multiple flags have probability $O(\varepsilon^2)$. The errors from these exceptions are $O(\varepsilon u + \varepsilon^2) = O(\varepsilon^{5/3})$.

Using $u^{3/2} = c_0\varepsilon + O(\varepsilon^{4/3})$ gives the following atom probabilities, with an $O(\varepsilon^{4/3})$ remainder in each entry:

Atom	P_ε	Q_ε
000	ε	$c_0\varepsilon$
Each of 011, 101, 110	σ	$\sigma + h\varepsilon$
Each of 001, 010, 100	0	$h\varepsilon$
111	$1 - 3\sigma - \varepsilon$	$1 - 3\sigma - (c_0 + 6h)\varepsilon$

Since $1 - c_0 = 7h$, the deviation at 000 is $-d_0\varepsilon + O(\varepsilon^{4/3})$ and at each of the other seven atoms it is $h\varepsilon + O(\varepsilon^{4/3})$. Its total variation is $d_0\varepsilon + O(\varepsilon^{4/3})$ and its Euclidean norm is $\sqrt{8/7}d_0\varepsilon + O(\varepsilon^{4/3})$, proving both upper bounds. $\square$

Corollary 5.15 follows from Theorem 5.14 and Proposition 5.13. The constant $1 - 2^{-3/2}$ is the fraction of the rare atom $P_\varepsilon(000) \approx \varepsilon$ that, by the Finner inequality, no compatible law with zero marginals near $\frac{1}{2}\varepsilon^{2/3}$ can retain.

6 Rejecting orders from distance, input size and parity

6.1 Distance-promised rejecting orders

The required order depends on the input law. We give bounds in terms of its distance from compatibility, its rational encoding length and its violation of a cubic parity certificate.

Corollary 6.1 (Distance-promised rejecting order). *Let G be a correlation scenario with K joint outcomes, using the private-source convention of Section 4.3, and let $\delta_2 = d_2(P, \mathcal{C}_G)$ and $\delta = d_{\text{TV}}(P, \mathcal{C}_G)$. If $P \in \mathcal{I}_n^{\text{NW}}(G)$ then*

$$\delta_2^2 \leq \frac{1 - \|P\|_2^2}{n}, \quad (37)$$

so for incompatible P

$$t_{\min}^{\text{NW}}(P) \leq \left\lfloor \frac{1 - \|P\|_2^2}{\delta_2^2} \right\rfloor + 1 \leq \left\lfloor \frac{K - 1}{4\delta^2} \right\rfloor + 1;$$

for the binary triangle, $t_{\min}^{\text{NW}}(P) \leq \lfloor 7/(4\delta^2) \rfloor + 1$.

Proof. The first inequality is (27), since some ω attains at most the mean and q_ω is compatible. Feasibility at order n then forces $n \leq (1 - \|P\|_2^2)/\delta_2^2$, and $\|P - Q\|_2^2 \geq \|P - Q\|_1^2/K = 4d_{\text{TV}}(P, Q)^2/K$ with $\|P\|_2^2 \geq 1/K$ gives the second form. The triangle has $K = 8$. $\square$

The proof uses no convexity of $\mathcal{C}_G$: we extract one compatible law from the witness. A numerical lower bound on δ therefore suffices to choose a rejecting order; the result does not supply that lower bound from the input law.

6.2 Rational input bit-length complexity

Proposition 6.2 (Required order in the input size). *Encode each of the eight rational atoms by its reduced numerator and positive denominator in ordinary binary, with delimiters of linear overhead. For a law P on three bits, let $\ell(P)$ be its total bit length. For each $H \in \{\text{NW}, \text{AI}, \text{exp}\}$,*

$$\max\{t_{\min}^H(P) : \ell(P) \leq B, P \notin \mathcal{C}_\Delta\} = 2^{\Theta(B)}.$$

Proof of Proposition 6.2. Lower bound. With $k = 2^j$, $\varepsilon = k^{-3}$ and $\sigma = 1/(2k^2)$, the eight atoms of P_ε are rationals with denominator $8k^9$ and numerators of $O(j)$ bits, so $\ell(P_\varepsilon) = O(j)$, while Proposition 5.13 gives $t_{\min}^H(P_\varepsilon) > k/2 = 2^{j-1}$. Choosing j proportional to B gives $2^{\Omega(B)}$.

Upper bound. By [20], $\mathcal{C}_\Delta$ is the image of a compact semialgebraic parameter space (six latent values per source and stochastic binary response tables) under a polynomial map, hence a closed semialgebraic subset of the simplex. The function $u(P) = d_2(P, \mathcal{C}_\Delta)^2$ is therefore semialgebraic, and its graph is described by a fixed quantifier-free formula $\Phi(P, u)$ in integer polynomials of fixed degrees, by real quantifier elimination [1]. For a rational P of bit length B , substituting P and clearing denominators turns the atoms of Φ into univariate integer polynomials in u of fixed degree and coefficient height $2^{O(B)}$. The set of u satisfying $\Phi(P, u)$ is the single point $u(P)$; discard identically zero specialized polynomials, whose comparisons have constant truth values. If none of the remaining polynomials vanished at $u(P)$, their signs would persist on a neighbourhood and the formula would define an interval. Thus some nonzero specialized polynomial f vanishes

at $u(P)$. Remove any factor u^r from f to obtain $g(u) = a_0 + \cdots + a_d u^d$ with $a_0 \neq 0$, the same nonzero root, and coefficient height at most $H = 2^{O(B)}$. For $0 < u < 1/(H+1)$,

$$\left| \sum_{i=1}^d a_i u^i \right| \leq H \sum_{i=1}^d u^i < \frac{Hu}{1-u} < 1 \leq |a_0|,$$

so $g(u) \neq 0$. Hence $u(P) \geq 2^{-O(B)}$ when P is incompatible, and Corollary 6.1 gives $t_{\min}^{\text{NW}}(P) \leq (1 - \|P\|_2^2)/u(P) + 1 = 2^{O(B)}$; the two smaller feasible sets reject no later. $\square$

With the binary triangle held fixed, a quantifier-free membership formula contains a bounded number of integer polynomials of bounded degree. Evaluating their signs at a rational input takes time polynomial in B . Quantifier elimination supplies such a formula in principle [1], although deriving and storing it may be impractical. For this fixed scenario, a polynomial-time decision procedure therefore exists alongside the exponential bound on the inflation order. For an incompatible rational input, one can also compute the first rejecting order by enumerating exact rational feasibility problems.

6.3 Rejecting orders from cubic parity certificates

Remark 6.3 (Cubic certificates and explicit rejecting orders). In sign notation, every $R \in \mathcal{C}_{\square}$ and every $R' \in \mathcal{C}_{\triangle}$ satisfy

$$\mathbb{E}_R A \mathbb{E}_R B \mathbb{E}_R [AB] + 8 R(ABCD = -1) \geq 0, \quad \mathbb{E}_{R'} A \mathbb{E}_{R'} B \mathbb{E}_{R'} C + 10 R'(ABC = -1) \geq 0.$$

Replacing C by $-C$ on the square, and A by $-A$ on the triangle, preserves the compatible sets and gives the companion certificates $\mathbb{E}_R A \mathbb{E}_R B \mathbb{E}_R [AB] + 8 R(ABCD = 1) \geq 0$ and $-\mathbb{E}_{R'} A \mathbb{E}_{R'} B \mathbb{E}_{R'} C + 10 R'(ABC = 1) \geq 0$.

To prove the certificates, use $|xyz - x'y'z'| \leq |x - x'| + |y - y'| + |z - z'|$ on $[-1, 1]^3$. On the square, fix z in the proof of Lemma 4.5: the moments $\mathbb{E}A$, $\mathbb{E}B$, $\mathbb{E}[AB]$ differ from those of the comparison pair (A'_z, B'_z) by at most $2\Pr(A \neq A'_z)$, $2\Pr(B \neq B'_z)$ and $2e_z$, in total $4e_z \leq 8\delta_z$, and the comparison product $(\mathbb{E}f_z)^2(\mathbb{E}g_z)^2(\mathbb{E}h_z)^2$ is nonnegative; averaging over z gives the constant 8. On the triangle, the proof of Lemma 3.19 bounds the error of $\mathbb{E}A$ against $\sigma\tau$ by $1 - W$ and the errors of $\mathbb{E}B$ and $\mathbb{E}C$ against $\sigma\nu$ and $\tau\nu$ by $4(1 - W)$ together, and $(\sigma\tau\nu)^2 \geq 0$; the total $5(1 - W)$ equals $10 R'(ABC = -1)$. The certificates add the errors of the parity-rigidity argument [2, 3] instead of taking their maximum.

Let $\mathcal{F}$ be one of these four left-hand sides, read as a function of the law, let l_1, l_2, l_3 be the three moments in its product and $\Sigma(P) = 3 - l_1^2 - l_2^2 - l_3^2 \leq 3$. If $P \in \mathcal{I}_n^{\text{NW}}(\square)$, respectively $P \in \mathcal{I}_n^{\text{NW}}(\triangle)$, then

$$\mathcal{F}(P) \geq -\frac{\Sigma(P)}{n}, \quad \text{so} \quad t_{\min}^H(P) \leq \left\lfloor \frac{\Sigma(P)}{-\mathcal{F}(P)} \right\rfloor + 1 \leq \left\lfloor \frac{3}{-\mathcal{F}(P)} \right\rfloor + 1$$

for $H \in \{\text{NW}, \text{AI}, \text{exp}\}$ whenever $\mathcal{F}(P) < 0$. To obtain the order bound, use the compatible laws q_ω of Theorem 4.8. They satisfy $\mathbb{E}_\Gamma q_\omega = P$. For the three sign functions χ_i defining the moments $l_i = \mathbb{E}_P \chi_i$, put $h_i = \mathbb{E}_{q_\omega} \chi_i - l_i$. Applying (26) to $\mu \mapsto (\mathbb{E}_\mu \chi_i - l_i)^2$ gives

$$\mathbb{E}_\Gamma h_i = 0, \quad \mathbb{E}_\Gamma h_i^2 \leq \frac{1 - l_i^2}{n}.$$

The product $g(l) = \pm l_1 l_2 l_3$ has zero diagonal second derivatives and mixed second derivatives of absolute value at most one on $[-1, 1]^3$. Taylor's formula along the segment from l to $l + h$ therefore gives

$$|g(l + h) - g(l) - \nabla g(l) \cdot h| \leq \sum_{i < j} |h_i h_j| \leq \sum_i h_i^2.$$

The parity-error term of $\mathcal{F}$ is linear in the law, so its mean is unchanged. Averaging $\mathcal{F}(q_\omega) \geq 0$ now yields $0 \leq \mathbb{E}_\Gamma \mathcal{F}(q_\omega) \leq \mathcal{F}(P) + \Sigma(P)/n$. This proves the bound for NW, and (1) transfers it to AI and exp.

Navascués and Wolfe prove that a polynomial inequality valid on the compatible set holds on $\mathcal{I}_n^{\text{NW}}$ up to an error of order $1/n$ [15, Theorem 1, eqs. (34) and (A5)]. Their explicit constant depends on the number of sources; the bound $\Sigma(P) \leq 3$ here is independent of it. Cubic triangle inequalities of the same form, in $\mathbb{E}A$, $\mathbb{E}B$, $\mathbb{E}C$ and $\mathbb{E}[ABC]$, appear among the order-two inequalities of Wolfe, Spekkens and Fritz [22, Section IV C, eqs. (51)–(54) of arXiv v5]; the certificates above are not among them. A certificate can reject a law that satisfies the Finner inequality (2) at every atom.

For example, the W law, uniform on the three outcomes with exactly one 1, has $-\mathbb{E}A\mathbb{E}B\mathbb{E}C + 10P(ABC = 1) = -1/27$ and $\Sigma = 8/3$, so the bound gives $t_{\min}^H \leq 73$. Its incompatibility with the triangle is Example 2 of [22], and that it satisfies the Finner inequality is noted in [19, footnote 4]. Navascués and Wolfe report that their hierarchy rejects it at order three [15, footnote 2], so the order above is far from tight for this law.

7 Open problems

Question 7.1 (Worst-case convergence exponent). For the square and the triangle, Corollary 4.9 gives $\Omega(1/t) \leq H_t^H \leq O(t^{-1/2})$ for each hierarchy H . What is the sharp asymptotic rate? Does a power-law exponent exist, and does the rate depend on the hierarchy?

For longer cycles Theorem 3.21 gives $\Omega(m^{-2}t^{-2})$, and for the five-vertex path Theorem 3.26 gives $\Omega(t^{-2})$.

Question 7.2 (Exact leading constant along the defect family). Does $\varepsilon^{1/3}t_{\min}^H(P_\varepsilon)$ converge? If so, what is its limit, and do the three hierarchies share it?

The current bounds give the interval $[\frac{1}{2}, 4 - 2\sqrt{3}]$.

Question 7.3 (Separation of the hierarchies). Is the factor $3/2$ of Theorem 4.10 sharp, and does a bounded order conversion hold on every pair-source scenario?

Proposition 4.11 gives finite-order separations. It remains open whether some fixed scenario satisfies $\mathcal{I}_t^{\text{AI}} \subsetneq \mathcal{I}_t^{\text{NW}}$ at every order $t \geq 2$.

Question 7.4 (Beyond pair-source scenarios). Which scenarios with sources shared by three or more observers admit a finite characterizing order?

Remark 3.9 shows that the observer graph alone does not decide this.

Question 7.5 (Explicit rejecting orders beyond Finner). Can one obtain computable rejecting-order bounds, with explicit constants, for every incompatible rational law on the triangle or square, including laws that satisfy all Finner and cubic parity inequalities considered here?

Theorem 5.11 covers the laws that violate a Finner inequality, and Remark 6.3 covers the laws on the square and the triangle that violate a cubic parity certificate. The Spiral-inflation inequalities of [22, Section IV C] and equation (31) of [15] are closed-form criteria that reject some Finner-satisfying laws at order two, but they give no rejecting order beyond a fixed level. The constants 8 and 10 of the certificates are not known to be sharp. The distance and bit-length bounds apply more broadly, but computing the distance or extracting the constants from quantifier elimination may be impractical.

Question 7.6 (Quantum and no-signalling analogues). For which quantum or no-signalling inflation hierarchies can one construct incompatible laws with an unbounded first rejecting order and determine its growth?

8 Verification and availability

The sources, exact certificates and Lean development are available at <https://github.com/williamjblair/inflation-termination>. The classification, witness constructions and asymptotic bounds have analytic proofs in the paper. The finite threshold claims in Proposition 4.11 also depend on rational primal and dual certificates, checked with integer and rational arithmetic and Sturm sequences. Figure 6 plots the certified endpoints. The triangle brackets at orders 11 to 13 quoted after Proposition 4.12 rest on certificates of the same kind. The exact checks of the cubic certificates of Remark 6.3 on laws of random finite models are falsification tests; we prove the certificates in the text, and the proofs do not depend on these checks. Appendix C describes the formats and finite coverage; the repository’s `paper/review/VERIFICATION_GUIDE.md` gives replay commands and checker details.

The Lean development proves the binary classification equivalence and existence of strictly positive witnesses for finite latent alphabets, including the structural lemmas, double-star reconstruction, cycle and five-path obstructions, their distance lower bounds and local flips. Rationality of the classification witnesses follows from the finite constructions in the paper; the mapped Lean statements do not assert it. The terminal declarations for the corrected square and triangle witnesses prove the case $q = 1/(16t)$ of Theorems 4.3 and 4.4, with the constants $c = 5$ and $c = 4$ in the density; the theorems in the paper cover $0 < q \leq (16/15)^{2/t} - 1$ and $0 < q \leq (9/8)^{2/t} - 1$ with $c = mR^{m-1}$. The passage to arbitrary latent spaces uses the cardinality theorem of [20], which is not formalized here. A separate triangle development proves the defect construction and nontermination for arbitrary latent probability spaces and measurable stochastic responses, the fan bounds, and the finite bounds on the defect family. It also proves the triangle norm estimates from Theorem 4.8.

The general convex-order inequality, the corrected densities beyond $q = 1/(16t)$, the max-moment estimate and survivor corollaries, order conversion, count-moment and threshold results including Proposition 4.12, the cubic certificates and rejecting orders of Remark 6.3, distance asymptotics, bit-length theorem and supplementary inequalities are not formalized. Neither are the larger-observed-alphabet transfer or the limiting statements for the defect family. Declaration-level coverage appears in `README.md` and `formalization.yaml`; the axiom audit lists 113 declarations.

The binary, finite-latent classification is registered as PALOMAR-2026-09-14-000009, version 1, at the Palomar registry. That registration concerns the stated classification theorem, not all results in this paper.

A Structural details of inflation

A.1 The root-sink lemma

We prove Lemma 3.4, using the following description of active trails.

Lemma A.1 (Trail criterion). *Let X, Y, Z be disjoint sets of copied observations in the order- t inflation of a pair-source scenario. Every trail between a member of X and a member of Y alternates copied observations and source copies, each source copy on it being a fork and each internal copied observation a collider. Such a trail is active given Z if and only if all its internal copied observations lie in Z . Consequently $X \perp_d Y \mid Z$ if and only if no trail from X to Y has all its internal copied observations in Z .*

Proof. Every edge of the inflated graph joins a source copy to a copied observation, so the nodes of a trail alternate between the two kinds. A source copy has no parents, so both of its neighbours on the trail are its children and it is a fork. It lies outside Z , since Z consists of copied observations, so it never blocks. An internal copied observation has no children, so both of its neighbours on the trail are its parents and it is a collider; a collider blocks unless it or one of its descendants lies in Z , and it has no descendants, so it blocks exactly when it is not in Z . A trail is active given Z when no node on it blocks, which is the stated condition. The last sentence is the definition of d -separation. $\square$

The shortest trails are the two-step ones, $x - S_e^m - y$, which have no internal observation and so are always active: two copied observations reading a common source copy are never d -separated.

Proof of Lemma 3.4. (i) If every component is injectable, then components with a common source copy would be joined, so distinct components are ancestrally independent and W is their union. Conversely, let $W = W_1 \cup \dots \cup W_m$ with the W_r injectable and pairwise ancestrally independent. Every edge of the shared-parent graph joins two members reading a common source copy, hence two members of the same W_r . So each component lies in one W_r , and a subset of an injectable set is injectable.

(ii) Necessity: injectable sets are expressible by the first rule, and if X and Y are ancestrally independent then by Lemma A.1 no trail from X to Y can avoid internal copied observations, and one with an internal copied observation is blocked by $\emptyset$; hence $X \perp_d Y \mid \emptyset$ and the join rule with $Z = \emptyset$ prescribes $p(x)p(y)$. Iterating gives the products of Definition 2.4.

Sufficiency: let Γ satisfy the prescriptions of Definition 2.4. We prove by induction on the generation of the expressible family that every generated set W is an AI set and that the law prescribed for it agrees with the marginal Γ_W . By (i) and the AI prescriptions, the marginal of an AI set is the product of $P_{\pi(K)}$ over the components K of its shared-parent graph.

For an injectable set the claim is the first AI prescription. For a marginal, the components of the shared-parent graph on a subset $W' \subseteq W$ are contained in those on W , hence injectable, so W' is again an AI set, and the marginal of the prescription is the prescription of the marginal.

Consider the join rule with $X \perp_d Y \mid Z$, where $X \cup Z$ and $Y \cup Z$ are already generated. Let K be a connected component of the shared-parent graph on $X \cup Y \cup Z$ and suppose K meets both X and Y . Among all paths inside K from a member of X to a member of Y choose one of minimal length. No internal vertex of it lies in X , since the subpath from that vertex to the Y -end would be shorter, and for the same reason none lies in Y ; so every internal vertex lies in Z . Inserting between consecutive vertices a source copy they share turns this path into a trail

from X to Y whose internal copied observations all lie in Z , which is active by Lemma A.1 and contradicts $X \perp_d Y \mid Z$. Hence every component of the shared-parent graph on $X \cup Y \cup Z$ lies in $X \cup Z$ or in $Y \cup Z$. Such a component is also a component of the shared-parent graph on the smaller set, since it is connected there and no member of the smaller set outside it is adjacent to it. As $X \cup Z$ and $Y \cup Z$ are AI sets, all these components are injectable, so $X \cup Y \cup Z$ is an AI set by (i).

For the prescribed law, let $\mathcal{K}$ be the set of components of the shared-parent graph on $X \cup Y \cup Z$, split as $\mathcal{X} \sqcup \mathcal{Y} \sqcup \mathcal{Z}$ according to whether a component meets X , meets Y , or is contained in Z ; by the previous paragraph no component meets both X and Y , so this is a partition. Write $\alpha = \prod_{K \in \mathcal{X}} P_{\pi(K)}$, $\beta = \prod_{K \in \mathcal{Y}} P_{\pi(K)}$ and $\mu = \prod_{K \in \mathcal{Z}} P_{\pi(K)}$, each evaluated at the coordinates of its component. The components of the shared-parent graph on $X \cup Z$ are the members of $\mathcal{X} \cup \mathcal{Z}$ together with the components of $K \cap Z$ for $K \in \mathcal{Y}$; write $\nu_{\mathcal{Y}}$ for the product of $P_{\pi(C)}$ over the latter, and $\nu_{\mathcal{X}}$ for the analogous product coming from $\mathcal{X}$ inside $Y \cup Z$. Then

$$\Gamma_{X \cup Z} = \alpha \mu \nu_{\mathcal{Y}}, \quad \Gamma_{Y \cup Z} = \beta \mu \nu_{\mathcal{X}}, \quad \Gamma_Z = \mu \nu_{\mathcal{X}} \nu_{\mathcal{Y}},$$

the last because the components of the shared-parent graph on Z are those of $K \cap Z$ for $K \in \mathcal{X} \cup \mathcal{Y}$ together with $\mathcal{Z}$. At a point where $\Gamma_Z > 0$ the factors μ , $\nu_{\mathcal{X}}$ and $\nu_{\mathcal{Y}}$ are positive and cancel,

$$\frac{\Gamma_{X \cup Z} \Gamma_{Y \cup Z}}{\Gamma_Z} = \alpha \beta \mu = \Gamma_{X \cup Y \cup Z},$$

which is the prescribed value. Where $\Gamma_Z = 0$ the rule prescribes 0, and $\Gamma_{X \cup Y \cup Z} \leq \Gamma_Z = 0$ there, so the two also agree on zero fibres. This completes the induction. $\square$

A.2 Induced-subgraph transport

Proof of Lemma 3.6. (i) Take a G -model for Q and keep the observations at $V(H)$. Since H is induced, every edge of G with both endpoints in $V(H)$ is an edge of H . An edge of G with exactly one endpoint v in $V(H)$ has its source read, among the retained observations, only by v ; append that source to the source of one edge of H at v , which exists because v is not isolated in H , and let the other endpoint of that edge ignore the appended component. The result is an H -model with the same observed law on $V(H)$.

(ii) Let Γ^H be a witness for P_H at order t and write $E_H(v)$ for the set of edges of H at v . Define a table Γ for G as follows: draw the copied observations $(O_v^j)_{v \in V(H), j \in [t]^{E_H(v)}}$ from Γ^H , draw independent fair bits ξ_w^k for $w \notin V(H)$ and $k \in [t]^{E(w)}$, and set

$$O_v^i := O_v^{i|E_H(v)} \quad (v \in V(H)), \quad O_w^k := \xi_w^k \quad (w \notin V(H)),$$

so that a copied observation at an H -vertex ignores the indices of the edges leaving H .

Symmetry: a permutation of the copy indices of an edge of H acts on the first family as the corresponding element of the symmetry group of the H -inflation and fixes the fair bits; a permutation of the copy indices of any other edge leaves every O_v^i with $v \in V(H)$ unchanged and permutes the fair bits among themselves, which preserves their law. **Diagonal law:** row r of the G -inflation restricted to $V(H)$ is row r of the H -inflation, and the outside entries of the t rows are $t|V \setminus V(H)|$ distinct fair bits, so the diagonal law is $P_H^{\otimes t} \otimes \text{Bern}(1/2)^{\otimes t(V \setminus V(H))} = P^{\otimes t}$.

Injectable prescriptions: let S be injectable in the G -inflation with $W = \pi(S)$. Its members at H -vertices, with the outside indices dropped, form an injectable set of the H -inflation, because

index erasure is still injective and the indices on the edges of H still agree; so their joint law under Γ is $(P_H)_{W \cap V(H)}$. Its members at outside vertices are distinct fair bits, independent of everything else, and $P_W = (P_H)_{W \cap V(H)} \otimes \text{Bern}(1/2)^{\otimes (W \setminus V(H))}$, as required. Ancestrally independent families: source copies of edges of H are source copies in G as well, so ancestrally independent injectable sets of the G -inflation have ancestrally independent H -parts, whose laws multiply under Γ^H , while the fair bits are independent of them and of each other. Hence Γ satisfies the conditions of Definition 2.4 whenever Γ^H does. The expressible case follows from the AI case by Lemma 3.4, applied in H and in G .

(iii) For $Q \in \mathcal{C}_G$, part (i) gives $Q_{V(H)} \in \mathcal{C}_H$, and marginalization does not increase total variation, so $d_{\text{TV}}(P, Q) \geq d_{\text{TV}}(P_{V(H)}, Q_{V(H)}) = d_{\text{TV}}(P_H, Q_{V(H)}) \geq d_{\text{TV}}(P_H, \mathcal{C}_H)$. Taking the infimum over Q gives the claim. $\square$

A.3 Injectable sets of the triangle

We justify the description of injectable sets used in Definition 2.4. In the order- t inflated triangle, a set S of copied observations is injectable when the map that erases copy indices is injective on S and, for any two members of S that are copies of variables sharing a source in the triangle, the copy index of that source agrees [22, Definition 4]. Injectivity allows at most one copy of each of A, B, C . If S contains A^{ij} and $B^{i'k}$ then $i = i'$; if it contains A^{ij} and $C^{j'k}$ then $j = j'$; if it contains B^{ik} and $C^{jk'}$ then $k = k'$. Hence $S \subseteq \Delta_{ijk}$ for suitable (i, j, k) , any missing index being chosen freely. Conversely every subset of a copied triangle is injectable. This proves the characterization for all t .

The degree- t diagonal condition (ii) of Definition 2.3 implies the lower-degree conditions of [15]: the law of any $g \leq t$ diagonal triangles is a marginal of $P^{\otimes t}$, hence $P^{\otimes g}$, and by symmetry the same holds for any g copied triangles with pairwise distinct indices in all three families.

B Order conversion and finite thresholds

B.1 Proof of square order conversion

Proof of Theorem 4.10. The right inclusion is Proposition 2.6 and the equality is Lemma 3.4. For the left inclusion let Γ be an order- n Navascués–Wolfe witness for P with $n = \lfloor 3t/2 \rfloor$ and let Γ' be its restriction to the copy indices $1, \dots, t$ of each source. Restriction preserves symmetry and the diagonal law of degree t , so Γ' witnesses $P \in \mathcal{I}_t^{\text{NW}}$; we check the AI prescriptions.

Fix an AI set of Γ' and decompose it into the connected components of its shared-parent graph, which are injectable and pairwise ancestrally independent. A component's ancestry uses two, three or four source types; the two-type components are single copied observations and correspond to edges of the four-cycle $X - Y - Z - W - X$ of source types. Any two subsets of a four-element set of sizes at least three, or of sizes three and two, intersect, so a component using at least three source types conflicts with every other component and needs a diagonal row of its own. Let b be the number of these.

The remaining components form a multigraph on the source-type four-cycle, with multiplicities n_A, n_B, n_C, n_D on the edges A, B, C, D . Opposite edges are disjoint, so the colours $\max(n_A, n_C)$ for the pair $\{A, C\}$ and $\max(n_B, n_D)$ for the pair $\{B, D\}$ suffice, and their sum equals the maximum source-type degree Δ : if $n_A \geq n_C$ and $n_B \geq n_D$ the sum is $n_A + n_B = \deg X$, and $\deg X$ dominates the other three degrees. So $b + \Delta$ rows suffice.

Choose a source type of residual degree Δ , say X , together with its two neighbours Y and W on the cycle. A component using at least three source types omits at most one, so it uses at least two of X, Y, W ; each of the Δ two-type components incident to X uses exactly two of them. The components have pairwise disjoint copied ancestries, and the three families X, Y, W supply $3t$ copies in Γ' , so $2b + 2\Delta \leq 3t$ and therefore $b + \Delta \leq \lfloor 3t/2 \rfloor = n$.

Assign each component to the diagonal row of its colour. Components sharing a row use disjoint source types, and distinct components use distinct copies within each family, so independent permutations of the copy indices of the four sources place all components on their assigned rows simultaneously. Under Γ the rows are independent with law P . Within one row the components have disjoint original source types, so Lemma 3.5, available because $n \geq 2$, makes P factor across them. Hence the joint law of the placed components is the product of their P -marginals. Undoing the permutations, the AI set of Γ' has exactly the prescribed product. $\square$

The factor $3/2$ comes from placing all components on diagonal rows at once; ancestrally disjoint collections of $\lfloor 3t/2 \rfloor$ injectable components whose source-type conflict graph is complete show that this method cannot do better.

B.2 Count-moment reductions

The reductions below replace feasibility for a parity-perfect target by a linear program on the counts of negative auxiliary signs. Let $m \in \{3, 4\}$ and use the notation of Section 3.4: at order t the copied observation O_v^{ij} reads the source copies $S_{v-1,i}$ and $S_{v,j}$. *Sign potentials* of height t are signs $s = (s_{g,i})$, $g \in \{0, \dots, m-1\}$, $i \in [t]$, one family for each source, and their *potential table* is $O_v^{ij} = s_{v-1,i} s_{v,j}$ as in (12). The degree tuple $k \in \{0, \dots, t\}^m$ of a set of potential signs counts its members in each family, and $S = \sum_g k_g$. Put

$$D_{\text{AI}}(t) = \left\{ k \in \{0, \dots, t\}^m : S \text{ even, } 2 \max_g k_g \leq S \right\},$$

$$D_{\text{NW}}(t) = \{ b_1 + \dots + b_t : b_r \in \{0, 1\}^m \text{ of even weight} \}.$$

A nonzero b_r of even weight has weight at least two, so summing $2 \max_g b_{r,g} \leq |b_r|$ over r gives $D_{\text{NW}}(t) \subseteq D_{\text{AI}}(t)$.

Lemma B.1 (Sign potentials). *Let $m \in \{3, 4\}$, $t \geq 1$ and $q \geq 0$ with $P_{m,q}$ a law.*

- (i) *If an assignment of signs to the copied observations of order t has product 1 on every copied cycle $\{O_v^{i_{v-1}i_v} : 0 \leq v < m\}$, $i \in [t]^m$, then it is the potential table of some s , and s is unique up to the global flip $s \mapsto -s$.*
- (ii) *Let ρ be a law on sign potentials of height t that is invariant under permutations of the indices within each family, and let D be $D_{\text{AI}}(t)$ or $D_{\text{NW}}(t)$. Suppose every product of potential signs whose degree tuple lies in D has ρ -mean $(-q)^{S/2}$. Then the law of the potential table is a witness placing $P_{m,q}$ in $\mathcal{I}_t^{\text{AI}}(C_m) = \mathcal{I}_t^{\text{exp}}(C_m)$ if $D = D_{\text{AI}}(t)$, and in $\mathcal{I}_t^{\text{NW}}(C_m)$ if $D = D_{\text{NW}}(t)$.*

Proof. (i) Two copied cycles that differ only in the index i_g share every observation except $O_g^{i_{g-1}i_g}$ and $O_{g+1}^{i_g i_{g+1}}$. Comparing the cycles with $i_g = i$ and $i_g = 1$ gives $O_g^{ai} O_g^{a1} = O_{g+1}^{ib} O_{g+1}^{1b}$ for all a, b , so both sides equal a sign $\lambda_{g,i}$ that depends on neither a nor b , and $\lambda_{g,1} = 1$. Applying this to the two indices of O_g gives $O_g^{ai} = \lambda_{g-1,a} O_g^{1i} = \lambda_{g-1,a} \lambda_{g,i} O_g^{11}$. Choose signs c_g with $c_{g-1} c_g = O_g^{11}$ for every g ; going around the cycle this is possible exactly because $\prod_g O_g^{11} = 1$,

which is the constraint at $i = (1, \dots, 1)$, and the choice is unique up to a global sign. Then $s_{g,i} = c_g \lambda_{g,i}$ are potentials for the assignment. If s and s' are potentials for the same assignment, then $s_{v-1,i} s'_{v-1,i} = s_{v,j} s'_{v,j}$ for all v, i, j , so this sign is the same for every family and every index, and $s' = \pm s$.

(ii) A permutation of the copy indices of source g permutes family g , so the law of the table is invariant. For a set F of copied observations, $\prod_{O \in F} O$ is the product of the potential signs in ∂F , the signs occurring an odd number of times. Let $F_1, \dots, F_r$ be pairwise ancestrally independent injectable sets, $U_i \subseteq F_i$ and $U = \bigcup_i U_i$. By Lemma 3.16, ∂U is the disjoint union of the sets ∂U_i , each of even size with at most one sign in each family. A nonempty ∂U_i has at least two elements, so it contributes at most $|\partial U_i|/2$ to each k_g ; hence the degree tuple of ∂U lies in $D_{\text{AI}}(t)$. If $D = D_{\text{AI}}(t)$ the hypothesis gives $\mathbb{E} \prod_{O \in U} O = (-q)^{|\partial U|/2} = \prod_i (-q)^{|\partial U_i|/2}$, and by the last claim of Lemma 3.16 and (9) the right side is the value of the corresponding character under the product of the $P_{m,q}$ -marginals of the blocks. Characters determine a law on a finite sign cube, so every injectable marginal and every ancestral-independence product is as prescribed; the diagonal rows are such blocks, and Lemma 3.4 gives the recursive prescriptions. If $D = D_{\text{NW}}(t)$, the diagonal law is determined by the characters $\prod_r \prod_{v \in F_r} O_v^{rr}$ with $F_r \subseteq \{0, \dots, m-1\}$. Distinct rows use distinct indices, so the boundary of such a product has degree tuple $\sum_r \mathbf{1}_{\partial F_r}$, where ∂F_r is the boundary of F_r in C_m , an edge set of even size, and the tuple lies in $D_{\text{NW}}(t)$. Its mean is $\prod_r (-q)^{|\partial F_r|/2}$, the value under $P_{m,q}^{\otimes t}$ by (9). $\square$

Lemma B.2 (Count-moment reduction). *Let $0 \leq q \leq 3 - 2\sqrt{2}$ and $t \geq 1$. Then $P_q \in \mathcal{I}_t^{\text{AI}}(\square)$ if and only if there are $w_a \geq 0$, $a = (a_X, a_Y, a_Z, a_W) \in \{0, \dots, t\}^4$, with $\sum_a w_a = 1$ and*

$$\sum_a w_a \prod_g K_{k_g}^{(t)}(a_g) = (-q)^{S/2} \quad \text{for every } k \in D_{\text{AI}}(t), \quad (38)$$

where

$$K_k^{(t)}(a) = \binom{t}{k}^{-1} \sum_j (-1)^j \binom{a}{j} \binom{t-a}{k-j}$$

is the normalized Krawtchouk polynomial.

Proof. The value $K_k^{(t)}(a)$ is the mean of the product of k distinct signs drawn uniformly without replacement from t signs of which a are negative.

Let Γ witness $P_q \in \mathcal{I}_t^{\text{AI}}(\square)$. Every copied square is injectable, so it has law P_q and product 1 almost surely, and by Lemma B.1(i) almost every table is a potential table. Choose its potentials with an independent uniform global flip and then apply independent uniform permutations within each family; the law of the table does not change, and the resulting law ρ on potentials is invariant under both operations. Under ρ the mean of a product of potential signs depends only on its degree tuple k , and equals $\sum_a w_a \prod_g K_{k_g}^{(t)}(a_g)$, where w is the law of the counts a_g of negative signs. Each $k \in D_{\text{AI}}(t)$ is the degree tuple of the boundary of an AI set, constructed below, whose prescribed character is $(-q)^{S/2}$; this gives (38). Conversely, given w , draw $a \sim w$ and place a_g negative signs at uniformly random positions of family g , independently across families. The resulting law on potentials is invariant within families, its means are the left sides of (38), and Lemma B.1(ii) with $D = D_{\text{AI}}(t)$ gives the witness.

It remains to realize each $k \in D_{\text{AI}}(t)$. Regard k_g as the number of terminals in source family g , arranged around the source cycle $X - Y - Z - W - X$. A path between two distinct source families gives an injectable block whose boundary consists of its endpoints. We construct

paths using at most t copies of each family, with distinct copies assigned to distinct paths. Write $(a, b, c, d) = (k_X, k_Y, k_Z, k_W)$ and, after a cyclic relabelling, assume $a + c \geq b + d$. Put $r = (a + c - b - d)/2$. The balance condition gives $0 \leq r \leq \min(a, c)$. Pair r terminals of X with r terminals of Z . The remaining X, Z terminals number $b + d$ and can be paired with the Y, W terminals along adjacent edges, since each of X, Z is adjacent to each of Y, W . Route each of the r opposite pairs through either Y or W . There are $2t - b - d$ unused copies in those two families, and

$$r = \frac{a + c - b - d}{2} \leq t - \frac{b + d}{2} \leq 2t - b - d.$$

These copies cover the internal vertices. Assigning distinct copies at every terminal and internal vertex makes the resulting injectable blocks ancestrally independent, and their union has boundary with degree tuple k . $\square$

Lemma B.3 (The Navascués–Wolfe form). *Under the hypotheses of Lemma B.2, $P_q \in \mathcal{I}_t^{\text{NW}}(\square)$ if and only if (38) holds for every $k \in D_{\text{NW}}(t)$. The same reductions hold for the triangle family $\Pi(-q, -q, -q) = P_{3,q}$, $0 \leq q \leq 1/3$, with three families of signs.*

Proof. Every copied square is carried to the first diagonal row by independent copy-index permutations, so under a witness it has law P_q and product 1 almost surely. The proof of Lemma B.2 then applies with $D_{\text{NW}}(t)$ in place of $D_{\text{AI}}(t)$, once every $k \in D_{\text{NW}}(t)$ is the degree tuple of a prescribed diagonal character. The boundaries of vertex sets of a cycle are exactly its edge sets of even size, so $k = b_1 + \dots + b_t$ is the degree tuple of $\prod_r \prod_{v \in F_r} O_v^{rr}$ with $\partial F_r = b_r$. For the triangle the same proofs apply with $m = 3$. A tuple $(a, b, c) \in D_{\text{AI}}(t)$ is realized by $(a + b - c)/2$, $(b + c - a)/2$ and $(c + a - b)/2$ single-observation blocks between the respective source pairs; these counts are nonnegative integers and use exactly the prescribed number of copies. $\square$

Lemma B.4 (Monotonicity along the parity direction). *Let $H \in \{\text{NW}, \text{AI}, \text{exp}\}$, $G \in \{\triangle, \square\}$, and $0 \leq q' \leq q$ in the range where the parity target is a law. If the parity target with parameter q lies in $I_t^H(G)$, so does the one with parameter q' .*

Proof. If $q = 0$, then $q' = 0$ and there is nothing to prove. For $q > 0$, by Lemmas B.2 and B.3 we may take a witness to be a law on sign potentials. Multiply every potential sign by an independent ± 1 variable of mean $\rho = \sqrt{q'/q}$. A character whose boundary has S elements is multiplied by ρ^S , so its value $(-q)^{S/2}$ becomes $(-q')^{S/2}$, while symmetry and the block structure of every prescription are unchanged. Lemma 3.4 covers the recursively expressible test. $\square$

B.3 Certified threshold brackets

Proof of Proposition 4.11. For each order and hierarchy, the artifact provides a nonnegative rational count distribution at the lower endpoint and rational dual functionals whose expectation polynomials are positive on intervals covering $[q_{\text{hi}}, q_{\text{cap}}]$. The checker verifies every prescribed moment, dual nonpositivity on all count configurations, and the interval signs by Sturm sequences. Here $q_{\text{cap}} = 1/3$ for the triangle and $q_{\text{cap}} = 1715728752/10^{10} < 3 - 2\sqrt{2}$ for the square. For each two-sided bracket, monotonicity (Lemma B.4) extends rejection at q_{cap} to the remainder of the physical range, and extends the primal witness down to zero. The verification guide linked in Section 8 gives the files and commands.

t	square		triangle	
	q_t^{AI}	q_t^{NW}	q_t^{AI}	q_t^{NW}
1	$3 - 2\sqrt{2}$	$3 - 2\sqrt{2}$	1/3	1/3
2	0.1324743290	≥ 0.1715628752	0.2679491900	1/3
3	0.0950228950	0.0950228950	0.1999999980	0.1999999980
4	0.0835279580	0.0950228950	0.1730047060	0.1999999980
5	0.0657380020	0.0657380020	0.1428571400	0.1428571400
6	0.0603176600	0.0645356170	0.1283741760	0.1381856900
7	0.0502565110	0.0502565110	0.1092774840	0.1092774840
8	0.0470271490	0.0494541280	0.1011416960	0.1070388380
9	0.0406784540	0.0406784540	0.0888129680	0.0888129680
10	0.0384981990	0.0400943600	0.0833179940	0.0870307060

Table 1: Certified brackets for the parity thresholds. A decimal entry x means $x \leq q_t^H \leq x + 4 \cdot 10^{-9}$; the exact entries at order one and triangle NW order two follow from the endpoint laws and Proposition D.7; the entry ≥ 0.1715628752 is a lower bound only, the family ending at $3 - 2\sqrt{2} \approx 0.1715728753$.

Order one accepts every target law. Proposition D.7 gives the triangle NW order-two endpoint. The strict comparisons follow from the rational endpoints. On the square, the points $3/20$ and $1/10$ lie in the respective gaps; Theorem 4.10 supplies the inclusion $\mathcal{I}_3^{\text{NW}} \subseteq \mathcal{I}_2^{\text{AI}}$. $\square$

Some dual expectation polynomials have simple roots, including $2 - \sqrt{3}$, $1/5$ and $1/7$ in the triangle cases. A root yields an upper obstruction; an exact threshold identity would also require a primal witness at that root.

B.4 Order conversion along the parity direction

Proof of Proposition 4.12. On the square this is Theorem 4.10 for $s \geq 2$, since $\mathcal{I}_t^{\text{NW}}(\square) \subseteq \mathcal{I}_{\lfloor 3s/2 \rfloor}^{\text{NW}}(\square) \subseteq \mathcal{I}_s^{\text{AI}}(\square)$, and it is trivial for $s = 1$. On the triangle, $D_{\text{NW}}(t)$ is the set of degree sequences of multigraphs with at most t edges on the three source families, since the even-weight vectors of $\{0, 1\}^3$ are the zero vector and the three pairs. A tuple k is such a degree sequence exactly when $S = \sum_g k_g$ is even, $2 \max_g k_g \leq S$ and $S \leq 2t$. A tuple of $D_{\text{AI}}(s)$ has even S with $2 \max_g k_g \leq S \leq 3s$, hence $S \leq 2 \lfloor 3s/2 \rfloor \leq 2t$, so $D_{\text{AI}}(s) \subseteq D_{\text{NW}}(t)$. Now let Γ witness $P_{3,q} \in \mathcal{I}_t^{\text{NW}}(\triangle)$. By the proof of Lemma B.3, Γ yields a law on sign potentials of height t , invariant within each family, under which every product of potential signs with degree tuple in $D_{\text{NW}}(t)$ has mean $(-q)^{S/2}$. Keep the potential signs with indices in $[s]$. Their law is invariant under permutations of $[s]$ within each family, and every product of them with degree tuple in $D_{\text{AI}}(s) \subseteq D_{\text{NW}}(t)$ has mean $(-q)^{S/2}$. Lemma B.1(ii) with $D = D_{\text{AI}}(s)$ places $P_{3,q}$ in $\mathcal{I}_s^{\text{AI}}(\triangle)$. $\square$

C Certificate formats and coverage

The files below are under `artifact/certificates/` in the repository <https://github.com/williamjblair/inflation-termination>, with checkers in the certificate directories and under `artifact/verifiers/`, and the manifest of SHA-256 identities at `artifact/MANIFEST.json`.

Structural certificate at order t . A JSON object with the root law (kind, $\varepsilon_t = 1/(2t^3)$ and $1 - \varepsilon_t$ as exact fractions), the value r_t , the list of t^3 roots D_{ijk} , for each of the $3t^2$ outputs the list of roots on its line, the eight target atoms of P_t as exact fractions, and the claimed counts of diagonal equations (8^t) and symmetry generators ($3(t-1)$). The checker recomputes every claimed quantity from the definitions and compares exactly. It computes the copied-triangle marginals by enumerating the 2^{3t-2} assignments of the roots on the three lines through a cell, so $t = 3$, with 2^7 assignments per triangle and $8^3 = 512$ diagonal equations, takes negligible time.

Full order-two table. A JSON object listing the 12 observation names in a fixed order, the denominator $2^{32} = (2t^3)^{t^3}$ at $t = 2$, and the nonzero numerators $\Gamma_2(w)$ of (31) for 166 assignments $w \in \{0, 1\}^{12}$; unlisted assignments have numerator 0. The checker verifies $\sum_w \Gamma_2(w) = 1$, nonnegativity, $\Gamma_2(w) = \Gamma_2(\pi w)$ for the three adjacent transpositions π and all 4096 assignments, and the 64 equations $\sum_{w: w_{\text{diag}}=(x,y)} \Gamma_2(w) = P_2(x)P_2(y)$ for the diagonal coordinates $(A^{11}, B^{11}, C^{11}, A^{22}, B^{22}, C^{22})$.

Negative control. The same object with the numerators of the all-one and all-zero assignments changed by -1 and $+1$. Normalization and nonnegativity hold; the diagonal equation at $(000, 000)$ fails. The checker must reject it.

Promised-rejection calibration. An exact Farkas certificate that the rational law $\frac{1}{2}(\delta_{000} + \delta_{111})$ is infeasible at order two with margin $1/2$, which confirms rejection of $R_{1/2}$ from Proposition 5.12. This finite check does not certify the general convergence constants.

Fan certificates. A JSON list of records, one per k in the list 4, 5, 6, 8, 10, 16, 32, 64, 128, 256, 512, 1024, 4096, 65536, with $\varepsilon = k^{-3}$, $\sigma = 1/(2k^2)$, the eight exact atoms of P_ε , the values m, z and the Finner gap $z^2 - m^3$, the passing order $\lfloor 1 + k/2 \rfloor$ with the exact Bernoulli certificate $(t-1)\varepsilon \leq \sigma$, the first order at which $v_t = tz - m - \binom{t}{2}m^2$ is positive together with the exact rational value of v_t there, and the first order at which the second-moment inequality fails. An all-order certificate records the identity $1 + \frac{1}{2}S(S-1) - S = \frac{1}{2}(S-1)(S-2)$ and its nonnegativity on integers by cases, so that (36) needs no order-by-order enumeration; the checker still enumerates all assignments of the fan for small t and all count orbits for larger t , and verifies that the index permutations used in (35) extend to permutations of $[t]$. Negative controls perturb the atoms, the parameters and the claimed orders, and the checker must reject each of them.

Threshold extension. The directory `exponent/odd/` holds certificates in the format of the parity-threshold files for the triangle at orders 11, 12 and 13 in both hierarchies, and for the square at order 11 in the ancestral-independence hierarchy, with brackets for two intermediate degree systems of the triangle at orders four and six. Its checker reuses the threshold checker record by record, binds each file to exactly these records, and adds exact checks of the degree-set identities behind Proposition 4.12 for orders up to 16, the containment $D_{\text{AI}}(s) \subseteq D_{\text{NW}}(t)$ for $s, t \leq 12$, and the bracket comparisons quoted after that proposition.

Cubic certificates. The checker in `beyond-finner/` is exact and self-contained. It evaluates the certificates of Remark 6.3 on the exact laws of random finite deterministic models, a

falsification test rather than a proof, and checks the algebraic identities behind its bound $\mathcal{F}(P) \geq -\Sigma(P)/n$ by enumeration at small n .

C.1 Finite coverage

The fan bounds determine the first rejecting orders of P_ε exactly at $\varepsilon = k^{-3}$ for $k = 4, 6, 8, 10, 16$: they are 4, 5, 6, 7, 10. At $k = 32, 128, 1024, 65536$, the passing construction gives lower bounds 18, 66, 514, 32770, while the fan gives upper bounds 19, 70, 551, 35122.

Certificate set	Range	Checked properties
Defect cube	$t = 1, 2, 3$	Root probabilities, line supports, symmetry, copied-triangle marginals, diagonal products and Finner gaps.
Full defect table	$t = 2$	All 4096 assignments; 166 nonzero entries; 12,288 symmetry equations and 64 diagonal equations.
Fan inequalities	Fourteen rational targets	Passing bounds, fan and second-moment rejection bounds, the pointwise count identity, index embeddings and negative controls.
Classification witnesses	Cycles C_3, C_4, C_5 and path: $t \leq 2$; corrected square: $t \leq 2$; corrected triangle: $t \leq 3$	Auxiliary density, nonnegative observed law, symmetry, prescribed marginals and ancestral products, target separation and local flips where supplied.
Parity thresholds	Triangle and square, both tests, $1 \leq t \leq 10$	Primal moments, dual signs on every count configuration, interval positivity; brute-force AI degree enumeration for $t \leq 3$ and explicit realizations through $t = 10$.
Threshold extension	Triangle, both tests, $t = 11, 12, 13$; square AI, $t = 11$	As above, with identical odd-order brackets, NW-key duals and the order-12 separation; degree-set identities.
Cubic certificates	Sampled compatible laws	Certificate signs on exact model laws; identities behind the bound $\mathcal{F}(P) \geq -\Sigma(P)/n$.

Table 2: Finite certificate coverage. Positivity of an auxiliary density does not imply full support of its observed pushforward. The unfipped cycle witnesses have parity constraints.

D Supplementary inequalities and extraction results

We give three complements to the main proofs: a stronger fan inequality and its square specialization, bounds for independent-cell constructions, and two limitations of extraction from a supported table. Throughout, total variation is half the ℓ^1 distance.

D.1 A fan inequality retaining pair marginals

Proposition D.1 (Pair-retaining fan inequality). *Let $n \geq 2$ and let P be a law on three bits with $a = P_A(0)$, $b = P_B(0)$, $c = P_C(0)$, $d = P_{AB}(00)$, $g = P_{AC}(00)$ and $z = P(000)$. If $P \in \mathcal{I}_n^{\text{NW}}$ then*

$$d + g + (n - 2)z \leq a + \binom{n}{2}bc. \quad (39)$$

Proof. For bits $b_1, \dots, b_n, c_1, \dots, c_n$ put $U = \sum_k b_k$, $V = \sum_k c_k$, $W = \sum_k b_k c_k$. The pointwise inequality is

$$UV - W \geq \frac{2}{n}((n-2)W + U + V - n). \quad (40)$$

To see it, multiply by n and collect terms on one side as $R = nUV - (3n-4)W - 2U - 2V + 2n$. Since $W \leq \min(U, V)$ and its coefficient is negative, it suffices to treat $W = \min(U, V)$; assume $U \leq V$, so R becomes $nUV - (3n-2)U - 2V + 2n$. For $U = 0$ this is $2(n-V) \geq 0$; for $U = 1$ it is $(n-2)(V-1) \geq 0$; for $U \geq 2$ it increases in V , so its minimum is at $V = U$, where it equals $n(U-1)(U-2) \geq 0$.

Now take the fan of the proof of Theorem 5.11: $\mathbf{A} = \mathbf{1}[A^{11} = 0]$, $\mathbf{B}_k = \mathbf{1}[B^{1k} = 0]$, $\mathbf{C}_k = \mathbf{1}[C^{1k} = 0]$, and set $b_k = \mathbf{B}_k$, $c_k = \mathbf{C}_k$ in (40). Definition 2.3 gives $\mathbb{E}\mathbf{A} = a$, $\mathbb{E}[\mathbf{A}\mathbf{B}_k] = d$, $\mathbb{E}[\mathbf{A}\mathbf{C}_k] = g$ and $\mathbb{E}[\mathbf{A}\mathbf{B}_k\mathbf{C}_k] = z$, since a permutation of the Y -indices carries each of these sets into the diagonal triangle Δ_{111} , and $\mathbb{E}[\mathbf{B}_k\mathbf{C}_\ell] = bc$ for $k \neq \ell$ by (35). Hence $\mathbb{E}(UV - W) = n(n-1)bc$. Multiplying (40) by $\mathbf{A}$ and using $UV - W \geq 0$ gives

$$n(n-1)bc \geq \frac{2}{n}(n(n-2)z + nd + ng - na),$$

which rearranges to (39). $\square$

Since $d \geq z$ and $g \geq z$, (39) implies the first inequality of (33). We apply it to the square through a conditioning that costs no order.

Lemma D.2 (Same-order transfer). *Let $n \geq 2$ and $P \in \mathcal{I}_n^{\text{NW}}(\square)$. Then $P_{ABC|D=1} \in \mathcal{I}_n^{\text{NW}}(\triangle)$ whenever $P_D(1) > 0$, and in every case*

$$P_{ABD}(001) + P_{ACD}(001) + (n-2)P(0001) \leq P_{AD}(01) + \binom{n}{2}P_B(0)P_{CD}(01). \quad (41)$$

Proof. Condition an order- n square witness on $D^{11} = \dots = D^{nn} = 1$, an event of probability $P_D(1)^n > 0$, and retain the complete A, B, C arrays. Independent permutations of the X - and Y -indices, together with the simultaneous permutation of the Z - and W -indices, realize the three source symmetries of the triangle, and the conditional diagonal law is $(P_{ABC|D=1})^{\otimes n}$. Applying (39) to the conditional law, using the independence of B and D forced at order $n \geq 2$, and clearing the positive denominator gives (41). If $P_D(1) = 0$ every term of (41) containing $D = 1$ vanishes and the inequality is trivial. $\square$

D.2 Independent-cell constructions

Let U_{ijk} , $(i, j, k) \in [n]^3$, be independent with a common law μ on a finite alphabet, let each observed type $V \in \{A, B, C\}$ carry a permutation-invariant kernel g_V giving its probability of output one from the n roots on its line, and let the responses be independently randomized across copied observations. The resulting table is a full order- n witness, and with the profiles $f_V(s) = \mathbb{E}[g_V(s, U_2, \dots, U_n)]$ and $m_V = \sum_s \mu_s f_V(s)$ the target is $\sum_s \mu_s \otimes_V \text{Bern}(f_V(s))$. Write $\mathcal{L}_n^{(\kappa)}$ for the laws obtained this way from a root alphabet of size κ .

Lemma D.3 (Profile variance). *For every admissible line kernel, $\text{Var}_\mu(f_V) \leq m_V(1 - m_V)/n$.*

Proof. Write $f = f_V$, $m = m_V$, $h = f - m$, $v = \mathbb{E}h(U)^2$ and $T = \sum_{i=1}^n h(U_i)$, and let $g = g_V(U_1, \dots, U_n)$. Independence gives $\mathbb{E}T^2 = nv$, and symmetry of the kernel gives $\mathbb{E}[(g - m)T] = n\mathbb{E}[h(U_1)(g - m)] = nv$. Cauchy–Schwarz then gives $(nv)^2 \leq \text{Var}(g)nv$. For $v > 0$ divide by nv and use $0 \leq g \leq 1$ to get $nv \leq \text{Var}(g) \leq m(1 - m)$; the case $v = 0$ is immediate. $\square$

Proposition D.4 (Independent-cell ceilings). *Let $n \geq 2$ and $P \in \mathcal{L}_n^{(\kappa)}$ for some finite κ . Then*

$$d_{\text{TV}}(P, \mathcal{C}_\Delta) \leq \frac{1}{2n},$$

and if P violates the Finner inequality $P(abc)^2 \leq P_A(a)P_B(b)P_C(c)$ at some atom, then

$$d_{\text{TV}}(P, \mathcal{C}_\Delta) \leq \frac{2}{n(n-1)^2} \leq \frac{8}{n^3}.$$

Proof. For any three-bit law, replacing the pair (B, C) by an independent pair and keeping the conditional law of A produces a compatible law, so

$$d_{\text{TV}}(P, \mathcal{C}_\Delta) \leq 2 \min\{|\text{Cov}(A, B)|, |\text{Cov}(A, C)|, |\text{Cov}(B, C)|\},$$

the factor being the exact four-cell identity for a binary pair. Under the independent-cell hypothesis $\text{Cov}(V, W) = \mathbb{E}_\mu[(f_V - m_V)(f_W - m_W)]$, so Lemma D.3 and Cauchy–Schwarz give $|\text{Cov}(V, W)| \leq \sqrt{m_V(1 - m_V)m_W(1 - m_W)}/n \leq 1/(4n)$, which is the first bound. For the second, relabel the violating atom's events as ones and order the marginals $a \leq b \leq c$, and write $t_n(x) = x + (1 - x)/n$ and z for the atom probability. Lemma D.3 gives $\mathbb{E}f_V^2 \leq m_V t_n(m_V)$, so

$$z \leq \mathbb{E}f_A f_B \leq \sqrt{ab t_n(a) t_n(b)} \leq \sqrt{ab} t_n(c),$$

and $t_n(c) \leq \sqrt{c}$ whenever $c \geq (n - 1)^{-2}$, since with $y = \sqrt{c}$ one has $ny - 1 - (n - 1)y^2 = (1 - y)((n - 1)y - 1) \geq 0$. Hence $z^2 > abc$ forces $\max(a, b, c) < (n - 1)^{-2}$, and applying Lemma D.3 to the event indicators gives $|\text{Cov}(V, W)| \leq \sqrt{m_V m_W}/n < 1/(n(n - 1)^2)$. $\square$

D.3 Reweighting a supported table

For a deterministic table ω on the copied observations and index distributions u, v, w on $[n]$ let

$$q_\omega(u, v, w)(a, b, c) = \sum_{i,j,k} u_i v_j w_k \mathbf{1}[(A^{ij}, B^{ik}, C^{jk}) = (a, b, c)],$$

let $\mathcal{R}(\Gamma)$ be the set of all $q_\omega(u, v, w)$ with $\omega \in \text{supp } \Gamma$, and let $e_\Gamma(P) = \min_{Q \in \mathcal{R}(\Gamma)} d_{\text{TV}}(P, Q)$. Write $\mathcal{C}_{\text{par}} = \{\Pi(st, su, tu) : s, t, u \in [-1, 1]\}$ for the parity-perfect part of $\mathcal{C}_\Delta$, as characterized by Lemma 3.18.

Proposition D.5 (Finite catalogue). *Every law attainable by reweighting a given binary-triangle table is attainable using at most eight existing indices from each source, with no table entry changed. Consequently $\mathcal{R}(\Gamma)$ is a union of model images drawn from a catalogue of deterministic subtables whose size does not depend on n , crudely bounded by $\sum_{r,s,t \leq 8} 2^{rs+rt+st} \leq 2^{201}$.*

Proof. Fix two of the three index distributions. The output law is the convex combination $q = \sum_i u_i q_i$ of points of the seven-dimensional simplex of eight outcomes. If more than eight weights are positive, an affine dependence moves them without changing q until one weight vanishes; repeat, then compress the second and third source in turn. $\square$

Theorem D.6 (Identity transfer). *Let f be a polynomial of degree $d \leq n$ in the eight atom probabilities, and let Γ be an order- n symmetric witness for P . If $f(q_\omega(u, v, w)) = 0$ for every $\omega \in \text{supp } \Gamma$ and all u, v, w , then $f(P) = 0$.*

Proof. Homogenize f to degree d using $\sum_x p_x = 1$ and let F_f be its symmetric multilinear polarization, so $F_f(p, \dots, p) = f(p)$. For a fixed table the function $G_\omega(u, v, w) = F_f(q_\omega, \dots, q_\omega)$ is homogeneous of degree d in each of the three weight blocks and vanishes on the product of simplexes, hence, by homogeneity after normalizing positive weight vectors, on the product of positive cones; so it is the zero polynomial. Select d distinct indices in each source family and read off the coefficient of the corresponding square-free monomial: it is a sum of terms $F_f(\delta_{\Delta_1(\omega)}, \dots, \delta_{\Delta_d(\omega)})$ over copied triangles with pairwise distinct indices in every family, and it vanishes. Average over all ordered injections $[d] \hookrightarrow [n]$ in each family and then over $\omega \sim \Gamma$. By the symmetry and the degree- d diagonal condition of Definition 2.3 the selected triangles are jointly distributed as $P^{\otimes d}$, so multilinearity gives

$$0 = \mathbb{E} F_f(\delta_{\Delta_1}, \dots, \delta_{\Delta_d}) = F_f(P, \dots, P) = f(P). \quad \square$$

Compatibility of the target alone does not give exact extraction at low order.

Proposition D.7 (Order-two obstruction). *Every parity-even three-bit law passes the order-two symmetric test. For the uniform parity law U , which is compatible, there is an order-two witness Γ with $e_\Gamma(U) = 1/2$.*

Proof. Draw an ordered pair $(x, y) \sim P^{\otimes 2}$ of parity-even triples. If $x = y$, let ω_{xy} be the constant table with that value everywhere. If $x \neq y$, the two triples differ at exactly two observed nodes, and those two nodes share a source; activate only that source, letting its index 1 produce the entries of x and its index 2 those of y at the two differing nodes, with the third node constant and all other source indices ignored. This defines the whole twelve-bit table. Its diagonal is (x, y) ; swapping the active source's two indices exchanges ω_{xy} with ω_{yx} , whose weights $P(x)P(y)$ agree; swapping inactive indices changes nothing. So the mixture is a normalized, nonnegative, symmetric table with diagonal law $P^{\otimes 2}$, which proves the first claim.

For U the four constant tables and the twelve oriented two-outcome tables each carry probability $1/16$. Every reweighting of a supported table is supported on at most two of the four parity-even outcomes, so for its support D one has $d_{\text{TV}}(U, Q) \geq U(D^c) \geq 1/2$, with equality for a nonconstant table whose active source is fair. Yet $U = \Pi(0, 0, 0) \in \mathcal{C}_{\text{par}}$. $\square$

The coefficient-extraction step in the proof of Theorem D.6 does not preserve inequalities, even at degree two. Let n be even and take the table $A^{ij} = b_i$ with half the b_i equal to zero and half to one, and B, C constant. The polynomial $f(q) = (q_A(1) - \frac{1}{2})^2$ is nonnegative on every reweighting, while its degree-two polarization on distinct source indices is

$$\frac{1}{n(n-1)} \sum_{i \neq j} \left(b_i - \frac{1}{2}\right) \left(b_j - \frac{1}{2}\right) = -\frac{1}{4(n-1)} < 0,$$

because the centered sum vanishes while its sum of squares is $n/4$. This identifies the loss in one route to a C/n distance bound; it is not a counterexample to such a bound. Proposition D.7 shows that any universal constant C with $e_\Gamma(P) \leq C/n$ for all n, P and Γ must satisfy $C \geq 1$.

Acknowledgments

The author used AI systems in preparing the manuscript and verification code and takes responsibility for the content. Section 8 distinguishes the analytic proofs, finite certificate checks and Lean formalization.

References

- [1] Saugata Basu, Richard Pollack, and Marie-Françoise Roy. *Algorithms in Real Algebraic Geometry*, volume 10 of *Algorithms and Computation in Mathematics*. Springer, 2 edition, 2006. doi:10.1007/3-540-33099-2.
- [2] Sadra Boreiri, Antoine Girardin, Bora Ulu, Patryk Lipka-Bartosik, Nicolas Brunner, and Pavel Sekatski. Towards a minimal example of quantum nonlocality without inputs. *Physical Review A*, 107:062413, 2023. arXiv:2207.08532, doi:10.1103/PhysRevA.107.062413.
- [3] Sadra Boreiri, Bora Ulu, Nicolas Brunner, and Pavel Sekatski. Noise-robust proofs of quantum network nonlocality. *Quantum*, 9:1830, 2025. URL: <https://arxiv.org/abs/2311.02182>, arXiv:2311.02182, doi:10.22331/q-2025-08-27-1830.
- [4] Cyril Branciard, Denis Rosset, Nicolas Gisin, and Stefano Pironio. Bilocal versus nonbilocal correlations in entanglement-swapping experiments. *Physical Review A*, 85(3):032119, 2012. arXiv:1112.4502, doi:10.1103/PhysRevA.85.032119.
- [5] José Mário da Silva, Alejandro Pozas-Kerstjens, and Fernando Parisio. Local models and Bell inequalities for the minimal triangle network. *Physical Review A*, 112(3):L030403, 2025. URL: <https://arxiv.org/abs/2503.16654>, arXiv:2503.16654, doi:10.1103/cg9v-wv2t.
- [6] Helmut Finner. A generalization of Hölder’s inequality and some probability inequalities. *The Annals of Probability*, 20(4):1893–1901, 1992. URL: <https://projecteuclid.org/journals/annals-of-probability/volume-20/issue-4/A-Generalization-of-Holders-Inequality-and-Some-Probability-Inequalities/10.1214/aop/1176989534.full>, doi:10.1214/aop/1176989534.
- [7] Thomas C. Fraser. A combinatorial solution to causal compatibility. *Journal of Causal Inference*, 8(1):22–53, 2020. arXiv:1902.07091, doi:10.1515/jci-2019-0013.
- [8] Thomas C. Fraser and Elie Wolfe. Causal compatibility inequalities admitting quantum violations in the triangle structure. *Physical Review A*, 98:022113, 2018. arXiv:1709.06242, doi:10.1103/PhysRevA.98.022113.
- [9] Tobias Fritz. Beyond Bell’s theorem: correlation scenarios. *New Journal of Physics*, 14(10):103001, 2012. URL: <https://arxiv.org/abs/1206.5115>, arXiv:1206.5115, doi:10.1088/1367-2630/14/10/103001.
- [10] Victor Gitton. Outer approximations of classical multi-network correlations. *arXiv preprint*, 2022. arXiv:2202.04103.
- [11] Joe Henson, Raymond Lal, and Matthew F. Pusey. Theory-independent limits on correlations from generalized Bayesian networks. *New Journal of Physics*, 16:113043, 2014. arXiv:1405.2572, doi:10.1088/1367-2630/16/11/113043.
- [12] Wassily Hoeffding. Probability inequalities for sums of bounded random variables. *Journal of the American Statistical Association*, 58(301):13–30, 1963. doi:10.1080/01621459.1963.10500830.
- [13] Laurens T. Ligthart, Mariami Gachechiladze, and David Gross. A convergent inflation hierarchy for quantum causal structures. *Communications in Mathematical Physics*, 401(3):2673–2714, 2023. URL: <https://arxiv.org/abs/2110.14659>, arXiv:2110.14659, doi:10.1007/s00220-023-04697-7.

- [14] Florence Jessie MacWilliams and Neil James Alexander Sloane. *The Theory of Error-Correcting Codes*, volume 16 of *North-Holland Mathematical Library*. North-Holland, Amsterdam, 1977.
- [15] Miguel Navascués and Elie Wolfe. The inflation technique completely solves the causal compatibility problem. *Journal of Causal Inference*, 8(1):70–91, 2020. URL: <https://arxiv.org/abs/1707.06476>, arXiv:1707.06476, doi:10.1515/jci-2018-0008.
- [16] Alejandro Pozas-Kerstjens, Nicolas Gisin, and Marc-Olivier Renou. Proofs of network quantum nonlocality in continuous families of distributions. *Physical Review Letters*, 130:090201, 2023. arXiv:2203.16543, doi:10.1103/PhysRevLett.130.090201.
- [17] Marc-Olivier Renou and Salman Beigi. Network nonlocality via rigidity of token counting and color matching. *Physical Review A*, 105:022408, 2022. arXiv:2202.00905, doi:10.1103/PhysRevA.105.022408.
- [18] Marc-Olivier Renou and Salman Beigi. Nonlocality for generic networks. *Physical Review Letters*, 128:060401, 2022. arXiv:2011.02769, doi:10.1103/PhysRevLett.128.060401.
- [19] Marc-Olivier Renou, Yuyi Wang, Sadra Boreiri, Salman Beigi, Nicolas Gisin, and Nicolas Brunner. Limits on correlations in networks for quantum and no-signaling resources. *Physical Review Letters*, 123(7):070403, 2019. URL: <https://arxiv.org/abs/1901.08287>, arXiv:1901.08287, doi:10.1103/PhysRevLett.123.070403.
- [20] Denis Rosset, Nicolas Gisin, and Elie Wolfe. Universal bound on the cardinality of local hidden variables in networks. *Quantum Information and Computation*, 18(11–12):910–926, 2018. URL: <https://arxiv.org/abs/1709.00707>, arXiv:1709.00707, doi:10.26421/QIC18.11-12-2.
- [21] Elie Wolfe, Alejandro Pozas-Kerstjens, Matan Grinberg, Denis Rosset, Antonio Acín, and Miguel Navascués. Quantum inflation: A general approach to quantum causal compatibility. *Physical Review X*, 11(2):021043, 2021. URL: <https://arxiv.org/abs/1909.10519>, arXiv:1909.10519, doi:10.1103/PhysRevX.11.021043.
- [22] Elie Wolfe, Robert W. Spekkens, and Tobias Fritz. The inflation technique for causal inference with latent variables. *Journal of Causal Inference*, 7(2):20170020, 2019. URL: <https://arxiv.org/abs/1609.00672>, arXiv:1609.00672, doi:10.1515/jci-2017-0020.